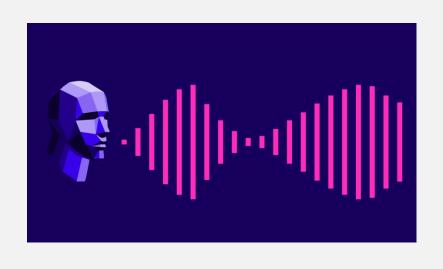
# Automatic Speech Recognition Research Master 2020-21



Radboud University

Nijmegen

Louis ten Bosch, E8.16





#### ASR is a dynamic field. Research is changing rapidly.

ASR is a broad field.
We only discuss the main issues.







#### Schedule

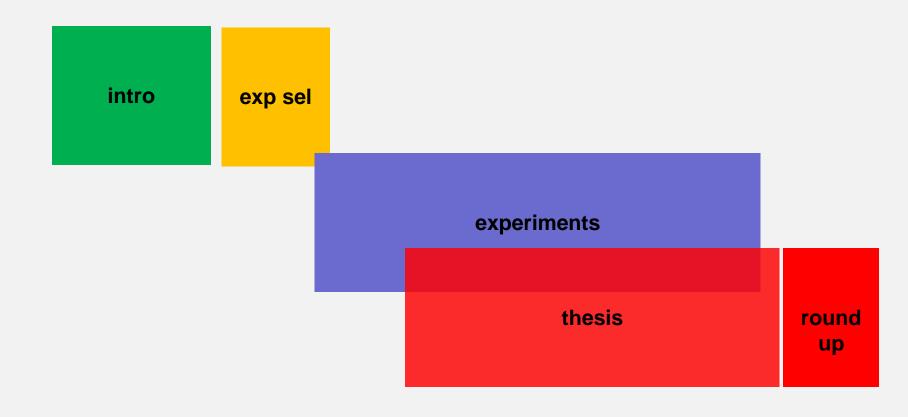
Wednesdays, start 15h30

- every week until June 2 (incl.)
  - with exceptions see course schedule





# Agenda

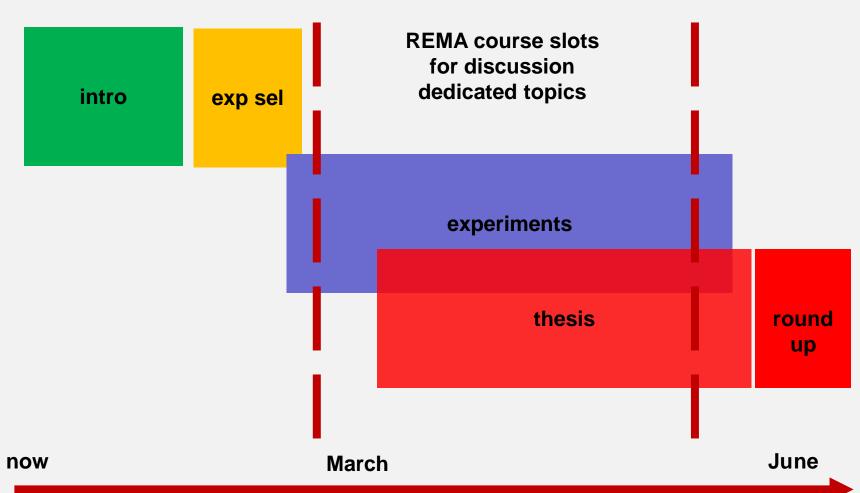


now June





# Agenda







# Global agenda this course

- Introductory part (approx. three weeks)
  - phonetics, speech, global ASR architecture
  - see also reading material (BrightSpace)
- Computational/experimental part (twelve weeks)
  - exploration, preparation (weeks 4 to 5)
  - experiments (individual or in groups; weeks 5 to 12)
  - reporting by individuals
- Goal: thesis (individual)
  - deadline: final version: June 15





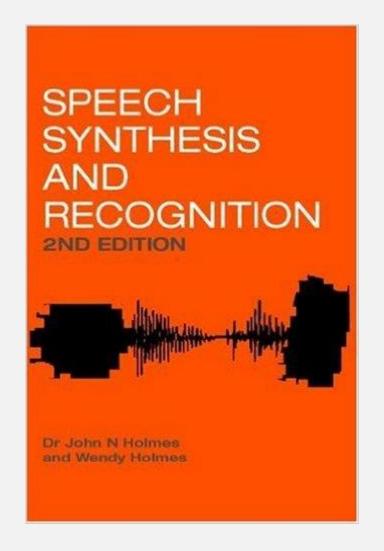
# Reading

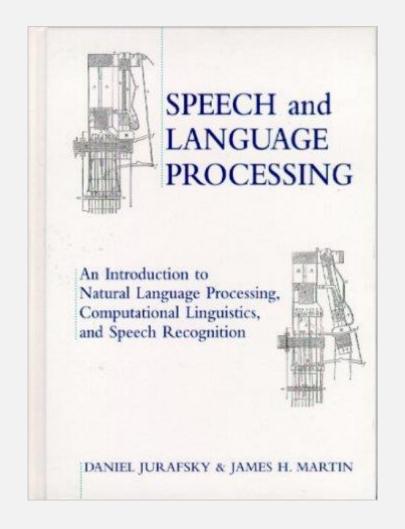
- Books: e.g.
  - Classical ASR: Holmes & Holmes (pdf on BrightSpace; see ASR chapters)
  - Jurafsky & Martin (more up-to-date)
- Papers (much more up to date)
  - there is a very large body of recent papers
    - about many highly specialized topics
  - Examples on BrightSpace
- These slides
  - will be updated on BrightSpace on a regular basis



#### Books











#### Aim of this course

- Provide background, insight
- Research oriented
- Indicate issues in ASR
- Invite future outlook

#### **Next Big Thing**

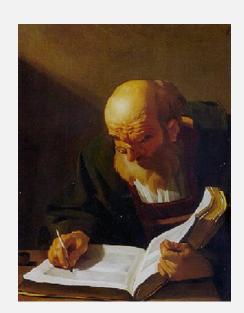






### **Thesis**

- Topic:
  - ASR is way too broad to be dealt with in 16 weeks
  - choose your topic narrow enough
  - usually based on an experiment
  - but may also be theoretical
- Starting point: a research question
- Experiment: individual or in a group
- Writing: individual
- Assessment: individual (see doc on BrightSpace)







# Experiments

- You may choose your own platform
- For experiments on Ponyland (science.ru.nl)
  - login
  - LINUX, perl/python, bash
- Proposals for experiments after the introductory stage (week 4 and later)



# Crowd sourcing speech



https://commonvoice.mozilla.org/nl



Het Common Voice-project is een initiatief van Mozilla om machines te helpen leren hoe echte mensen spreken. Spraak is natuurlijk, spraak is menselijk. Daarom zijn we enthousiast over het maken van een bruikbare spraaktechnologie voor onze machines. Maar voor het maken van spraaksystemen hebben ontwikkelaars een extreem grote hoeveelheid spraakgegevens nodig.

De meeste gegevens die door grote hedriiven worden gebruikt





# Options for experiments

#### just examples

- 1. ASR (audio → text) for command and control (Dutch) (github Jurriaan)
- 2. Keyword spotting for Dutch (Tara Sainath papers)
- 3. Personalized ASR
- 4. EARSHOT (Magnuson et al., 2020) (github Pepijn)
- 5. Estimation of #talkers in multi-talker audio input (audio  $\rightarrow$  {0,1,2,3,...})
- 6. Acoustic/phonological feature detectors (audio → vector)
- 7. Isomorphisms between latent structures in DNNs (X-H-Y, X-H2-Y)
- 8. Relation between *acoustic* and *symbolic* distance (log prob score and levenshtein distance, e.g. bull bill ball, morse horse force)
- 9. Automatic improvement readability of ASR output (+ punctuation)
- 10. Voice morphing and unmorphing (teams Alice, Bob, Marc)
- 11. Audio unshredding (continuity constraints)
- 12. Detection of fake audio (how)





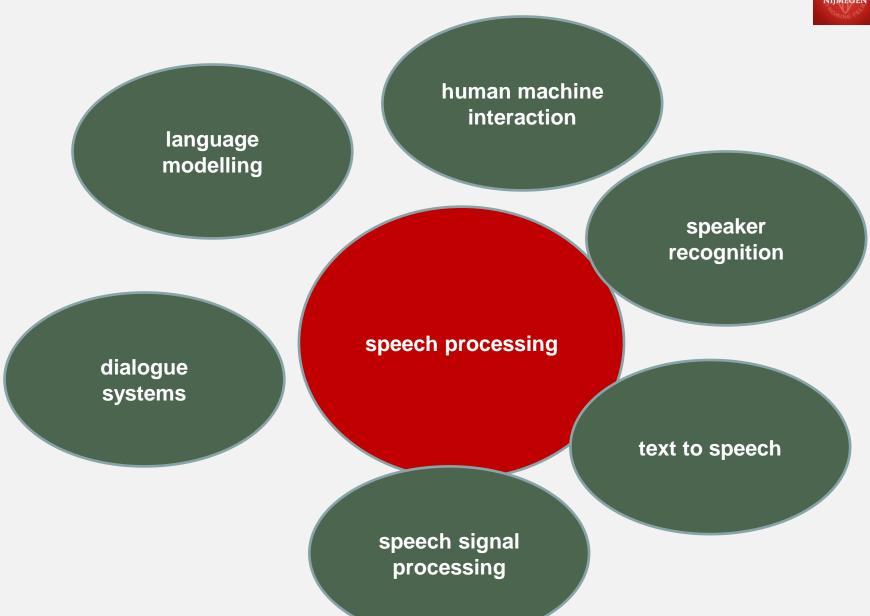
#### Audience?



- For now:
  - rondje
  - background, your own interest in ASR
  - AI/Data Science/Linguistics/...

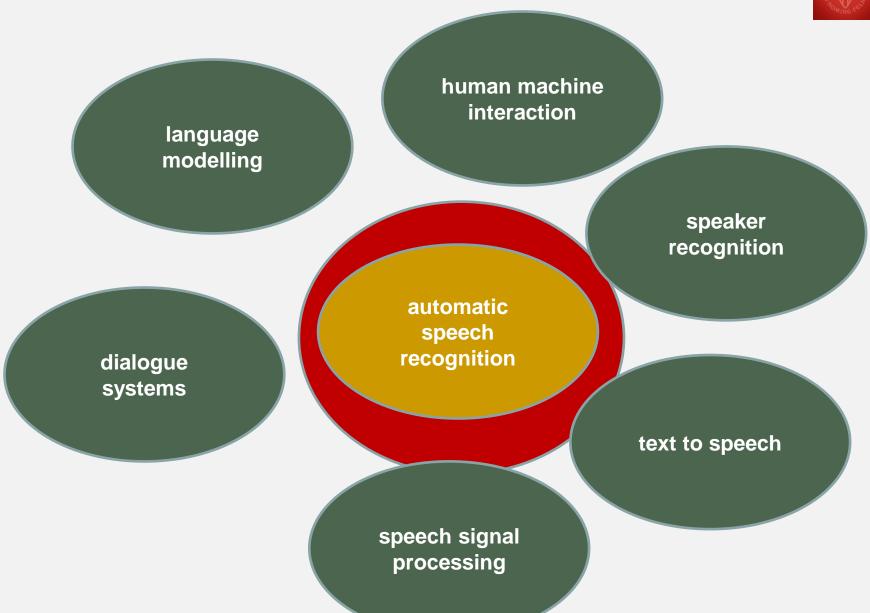
















# ASR = Speech to Text

Automatic Speech Recognition is the process in which a computer converts speech into text

No understanding. No semantics!



".... these cells indicate ..."

### ASR

# long history since 1750, breakthrough since DNNs (2010)



**Thomas Edison** invents the first dictation machine



1952

**IBM Shoebox** can understand 16 English words



1971



IBM Tangora, using the Hidden Markov Model, predicts upcoming phonemes in mobile devices 2006 speech





2011

1784







Bell Labs releases Audrey, capable of recognizing spoken digits with 90% accuracy - but only when spoken by its inventor





Harpy, created at Carnegie Mellon University, can comprehend 1,011 words - and some phrases



1986

The National Security Agency (NSA) starts using speech recognition to isolate key words in recorded speech



2008

Apple announces Siri, ushering in the age of the voiceenabled digital assistant



#### ASR is now common business

#### Human machine interaction

- speech is the most natural means of communication ~ 100000 y
- ease of communication
- safety (pilots, drivers, etc)
- communication with robots, care bots
- hands-free
- teacher, instructor, caregiver

#### Linguistics (e.g.), research

- automatic determination of dialect, speaking style, mood, ...
- large-scale processing of speech corpora





# ASR application areas

- Commercial speech recognition applications
  - voice dialing, ...
  - dictation programs (law, medical, office, ...; avoidance of RSI)
  - speech interface for games
  - embedded (robots, car, home, ambient intelligence...)
- Education/pathology (e.g.)
  - CAPT, CALL
  - assessment of pathological speech





# ASR application areas

- ASR =?= a model of Human Speech Processing (HSP)
  - By modelling human speech comprehension we may learn more about the comprehension process itself

#### Issues:

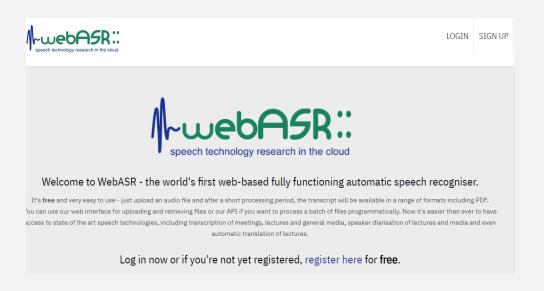
- plausibility of the processes
- plausibility of the representations
- how to get from subsymbolic information to symbolic information?





#### ASR as webservice

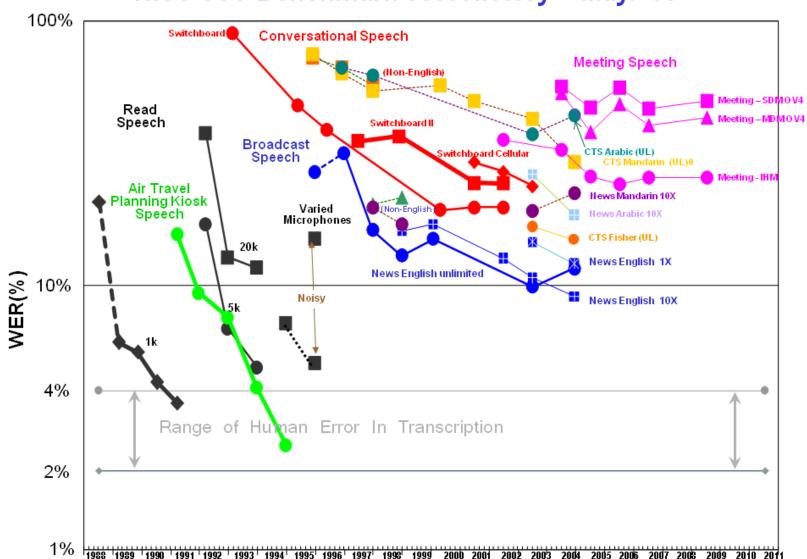
- webasr.org
- https://nlspraak.ewi.utwente.nl/ (ENG/NL)
- https://cls.ru.nl/online-asr/
- https://alphacephei.com/vosk/





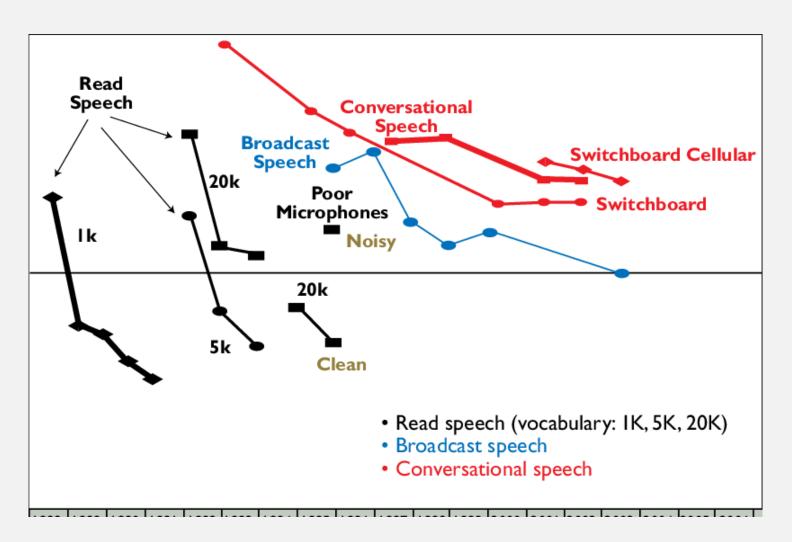


#### NIST STT Benchmark Test History - May. '09





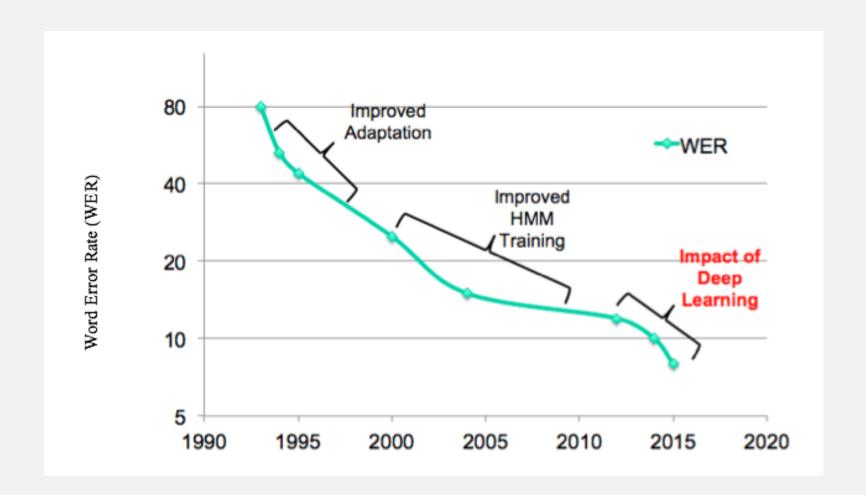




(By Xuedong Huang, 2016)

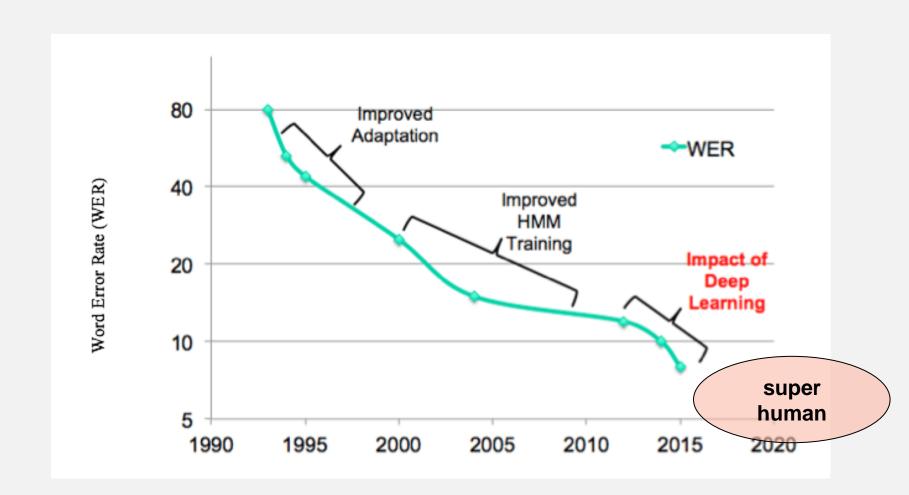
















# Four ways to look at ASR

- ASR as tool (black box) in larger system
  - In a dialogue system
  - As a tool for the hearing-impaired
  - In cases where typing is dangerous/impossible
- ASR as a computational model for human speech processing
  - Not straightforward, plausibility issues
- ASR as a tool to unravel the structure in the speech signal
  - Automatic segmentation into phones
- ASR as a research tool itself
  - WER minimization, noise robustness etc.

# ASR can be used in different ways



- For word recognition
  - open domain (any topic, most complex)
  - narrow domain, e.g., restaurants in Nijmegen (less complex)
  - Dictation (medical, court, police, broadcasting, subtitling, etc.)
- For keyword spotting
  - Simpler than full-blown speech recognition
  - For topic monitoring of thousands of telephone conversations in parallel (wakeup calls e.g. "hello siri", security, eavesdropping)
- For forced alignment
  - For research purposes (onsets and offsets of words, syllables)
  - Medical assessments of voice quality, etc.
- The main difference is in the WORD SEARCH SPACE





### Complexity of an ASR task

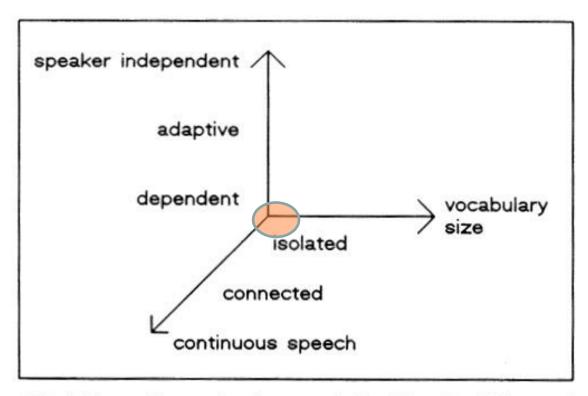


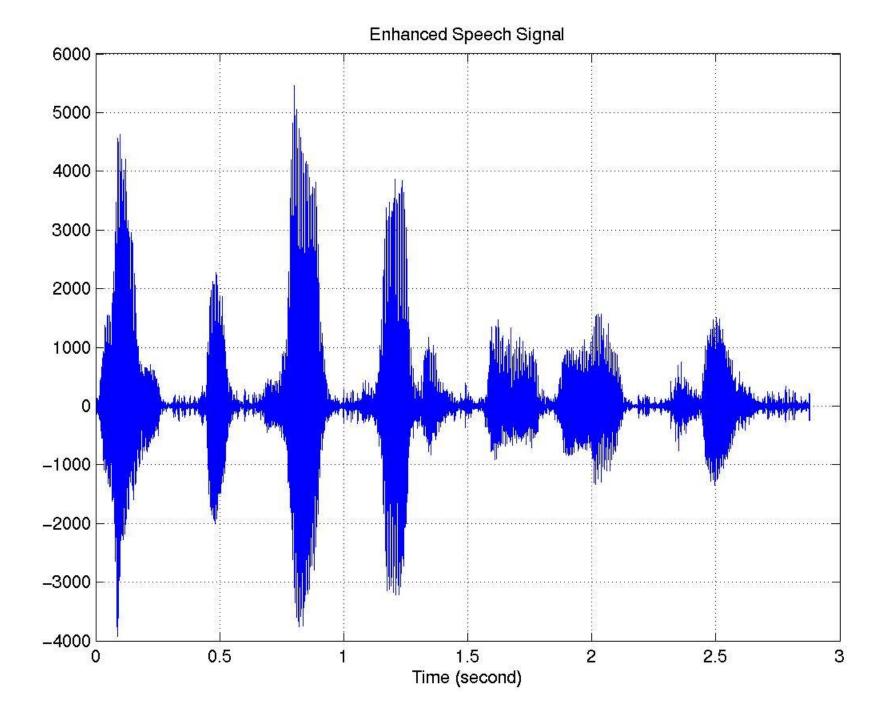
Figure 7.5. A three-dimensional space defined by the different functionalities provided by a recognizer.

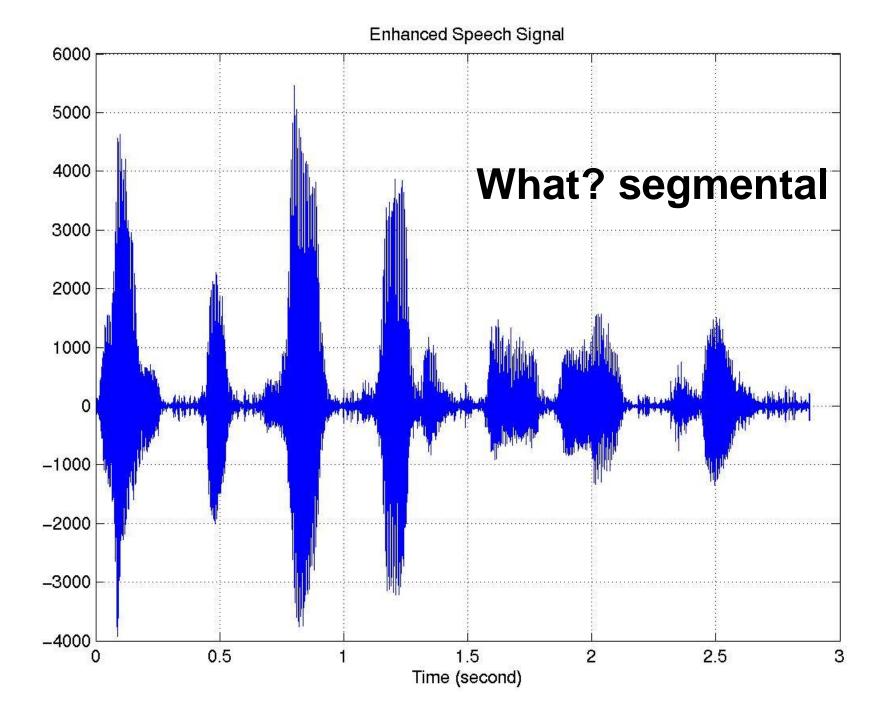


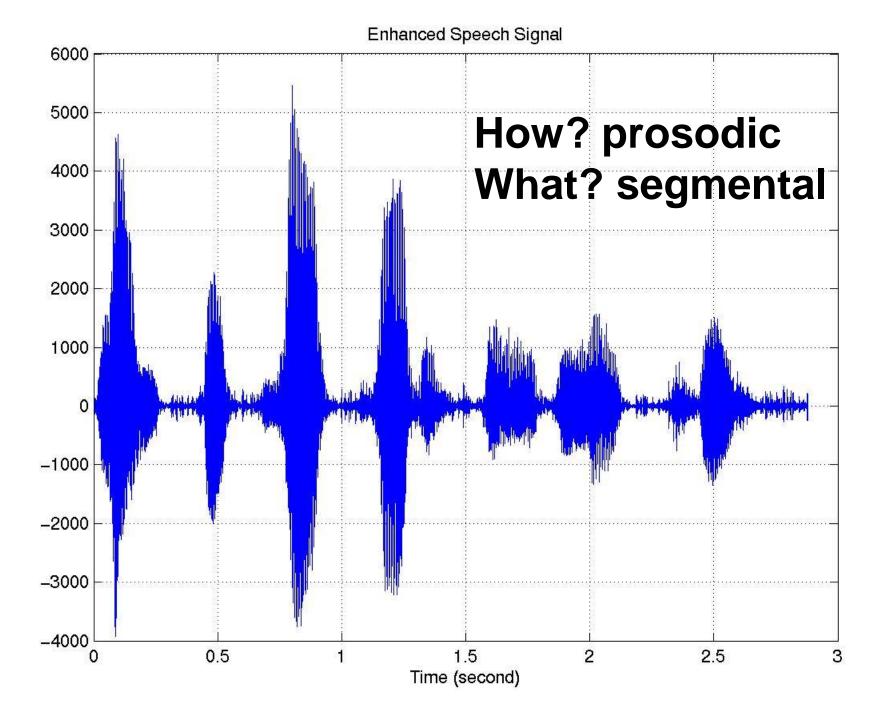


# Characteristics speech signal

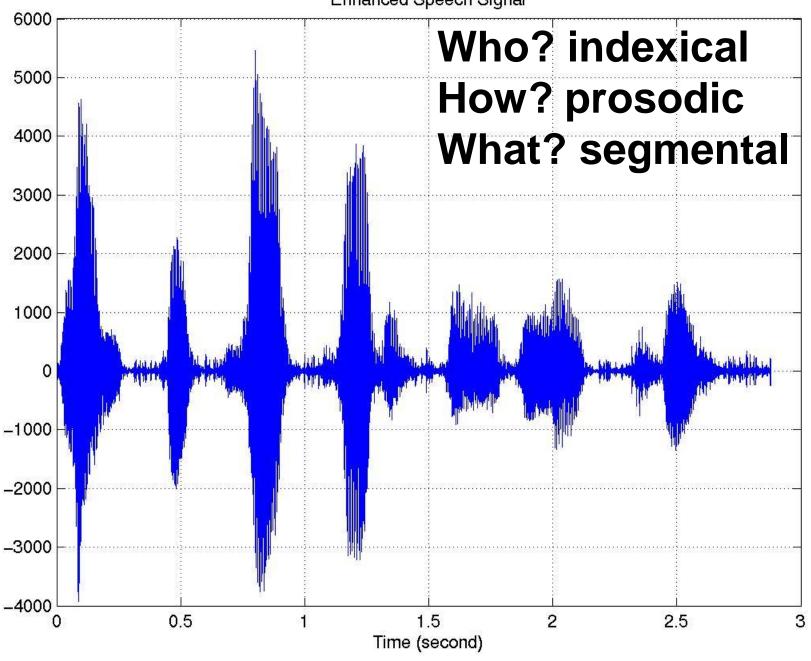
- No symbols
- No segmentation / boundaries / word boundaries
- Enormous variation (speakers, accents, mood, ...)
- Background noise
- Totally different from text!
- Speech is crucial for language evolution and language acquisition



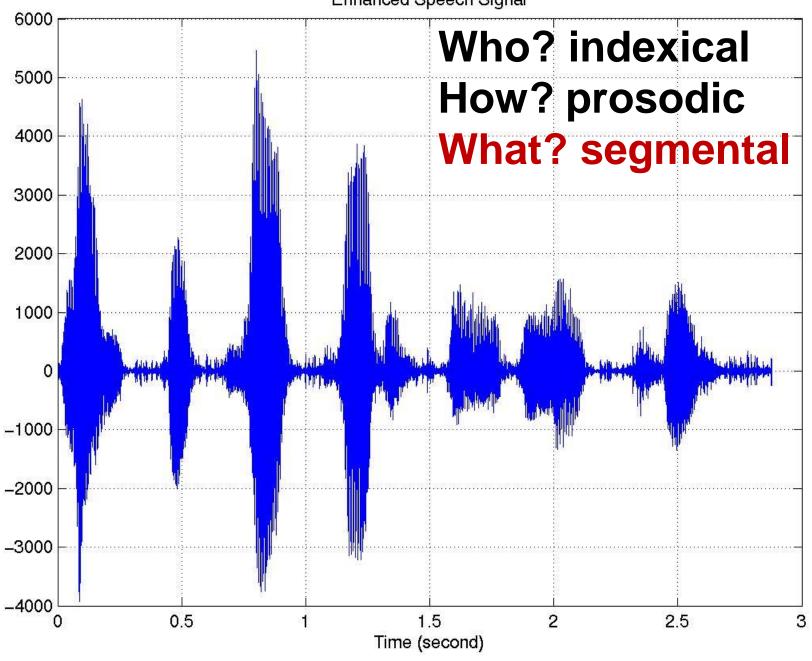








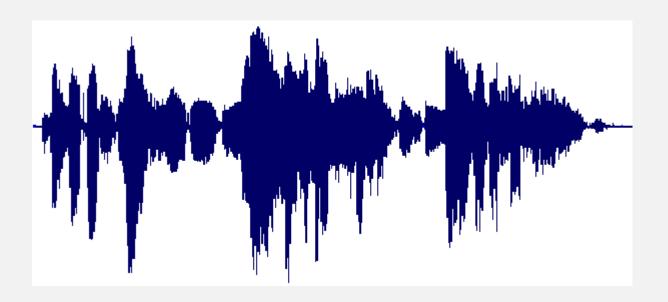
Enhanced Speech Signal







# from signal to words

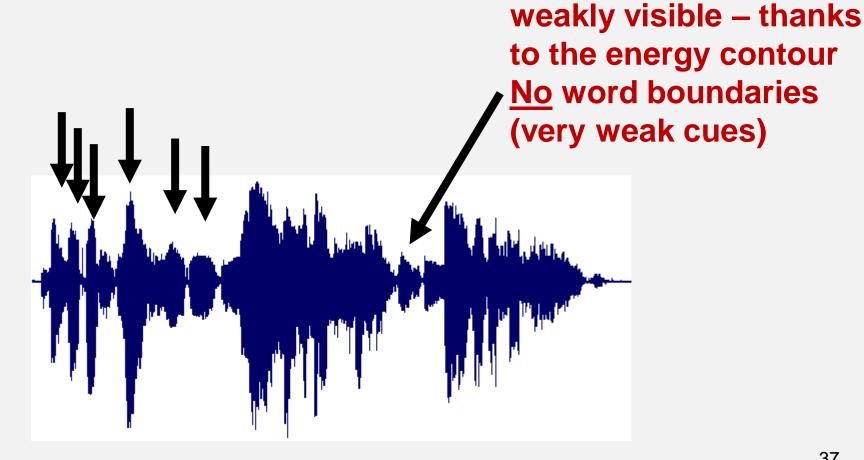






Syllable structure

#### from signal to words





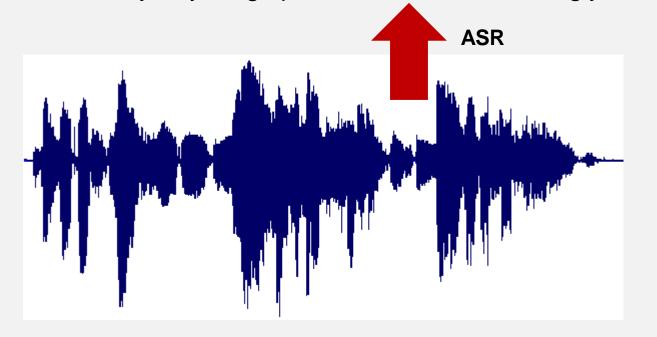


## from signal to words

There is no X: say("I", X) & speculative(X) ...

NLP/NLU

... I did not say anything speculative, but interestingly ...



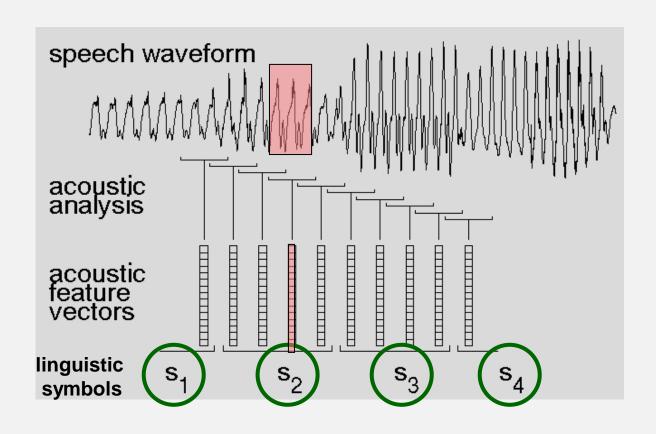
# High level description







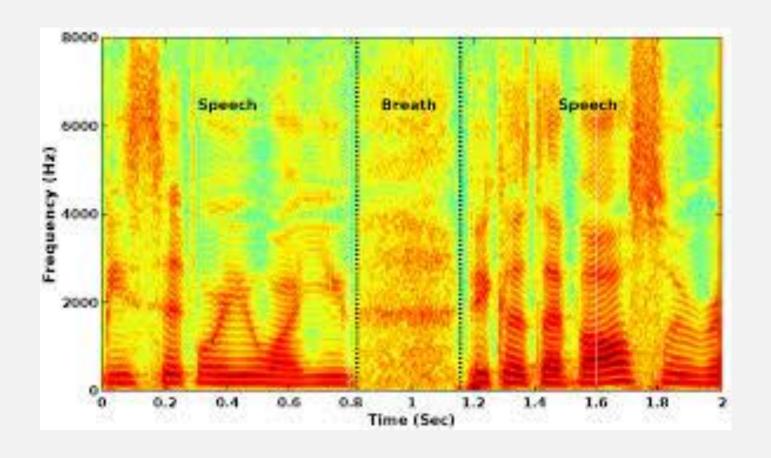
#### ASR: features ("front end")





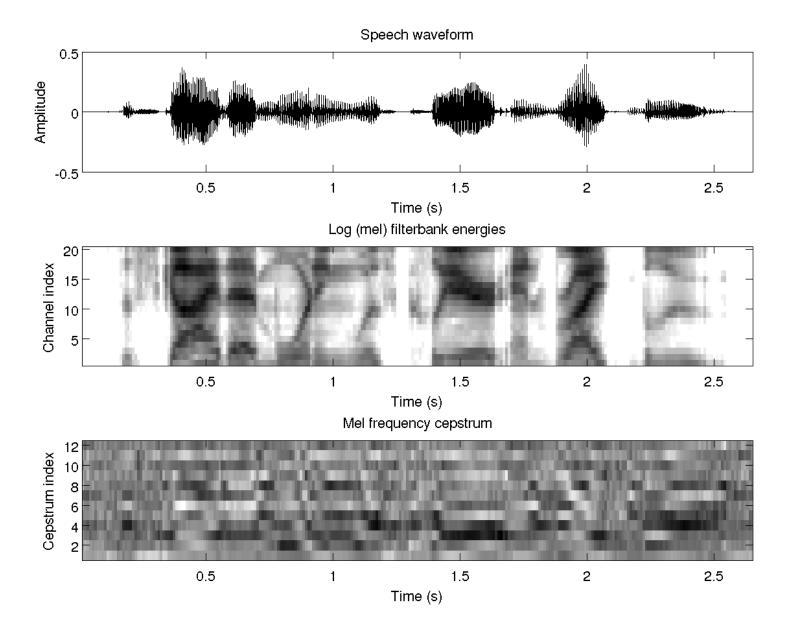
#### spectral analysis





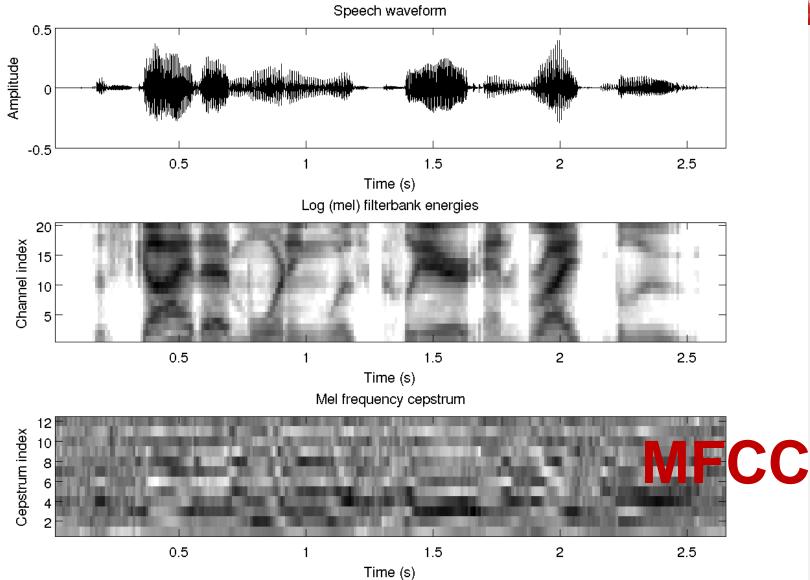














# MFCC vectors as audio representation



Very many sites show information about MFCCs, often with useful Python function calls

See e.g.

https://www.kaggle.com/ilyamich/mfccimplementation-and-tutorial

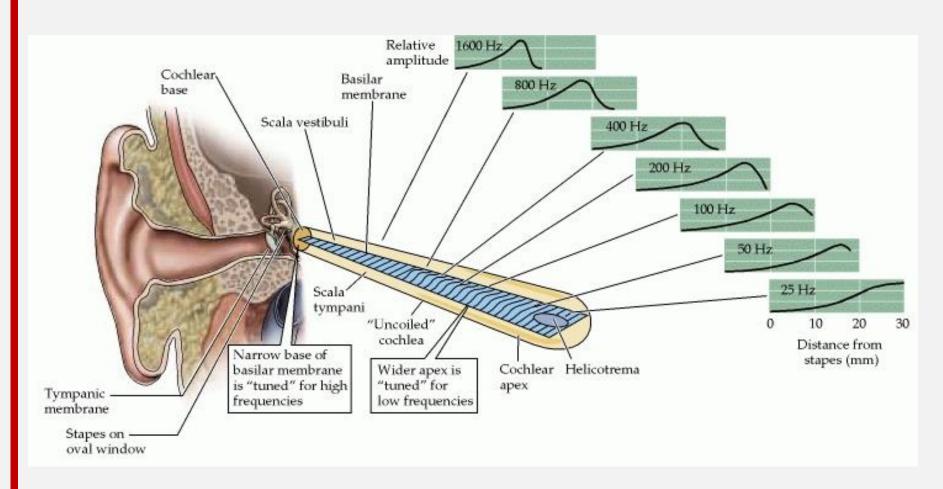
https://pypi.org/project/python\_speech\_features/

Librosa python library

https://librosa.org/doc/latest/index.html







Weber's law <a href="https://www.youtube.com/watch?v=hHG8io5qIU8">https://www.youtube.com/watch?v=hHG8io5qIU8</a>
Physics
Perception

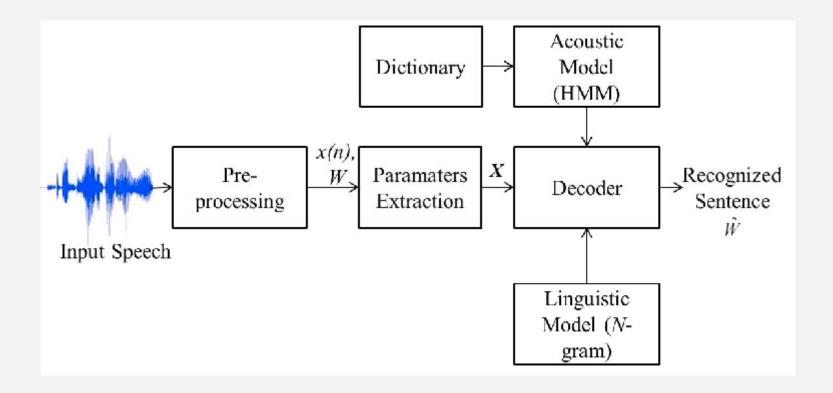
**Energy as function of frequency** 

log(E) as function of log(f)





#### ASR: classical architecture







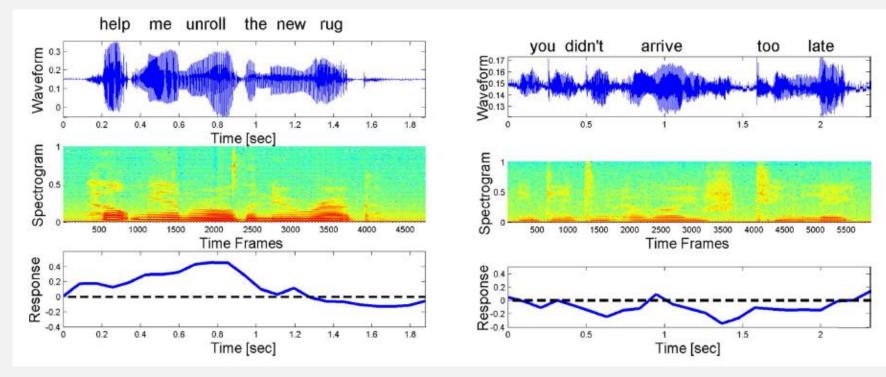
# "Free" recognition

```
audio 0.18 0.15 ik 1.00
audio 0.33 0.09 heb 1.00
audio 0.42 0.09 wel 0.95
audio 0.51 0.09 eens 0.88
audio 0.6 0.12 met 1.00
audio 0.72 0.39 serie 0.97
                                   Siri
audio 1.11 0.63 gesproken 1.00
audio 7.35 0.21 dit 0.87
audio 7.56 0.15 gaat 0.94
audio 7.71 0.12 wel 0.94
audio 7.83 0.30 enkel 0.56
audio 8.14 0.44 soepel 1.00
audio 24.99 0.09 dan 0.91
audio 25.08 0.12 zou 0.95
```





# **Keyword Spotting**





Discriminative Keyword Spotting for limited-data applications



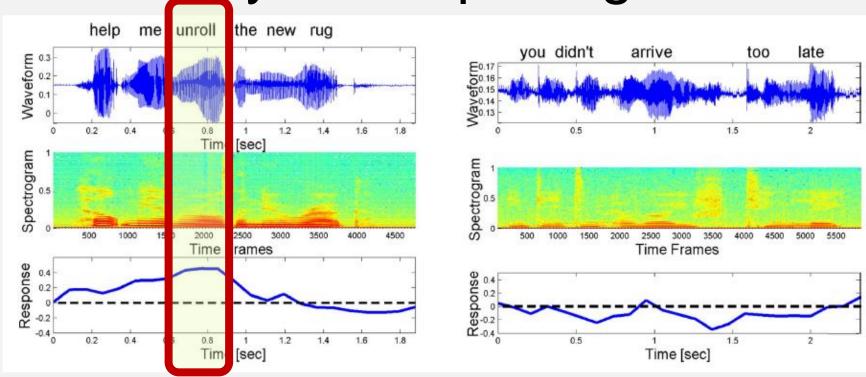
And rew and Erna Viterbi Department of Electrical Engineering, Technion-Israel, Institute of Technology, Haifa~32000, Israel.







# Keyword Spotting





Discriminative Keyword Spotting for limited-data applications



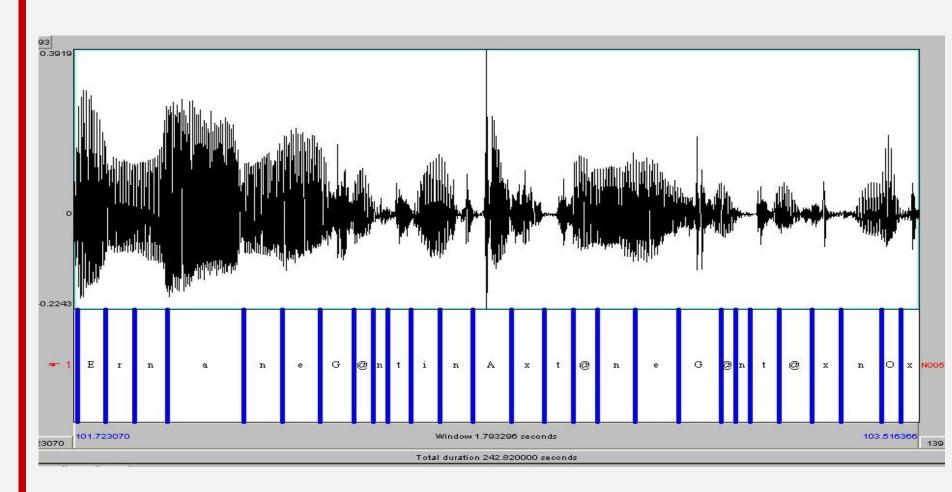
Hadas Benisty\*, Itamar Katz, Koby Crammer, David Malah

Andrew and Erna Viterbi Department of Electrical Engineering, Technion - Israel, Institute of Technology, Haifa 32000, Israel

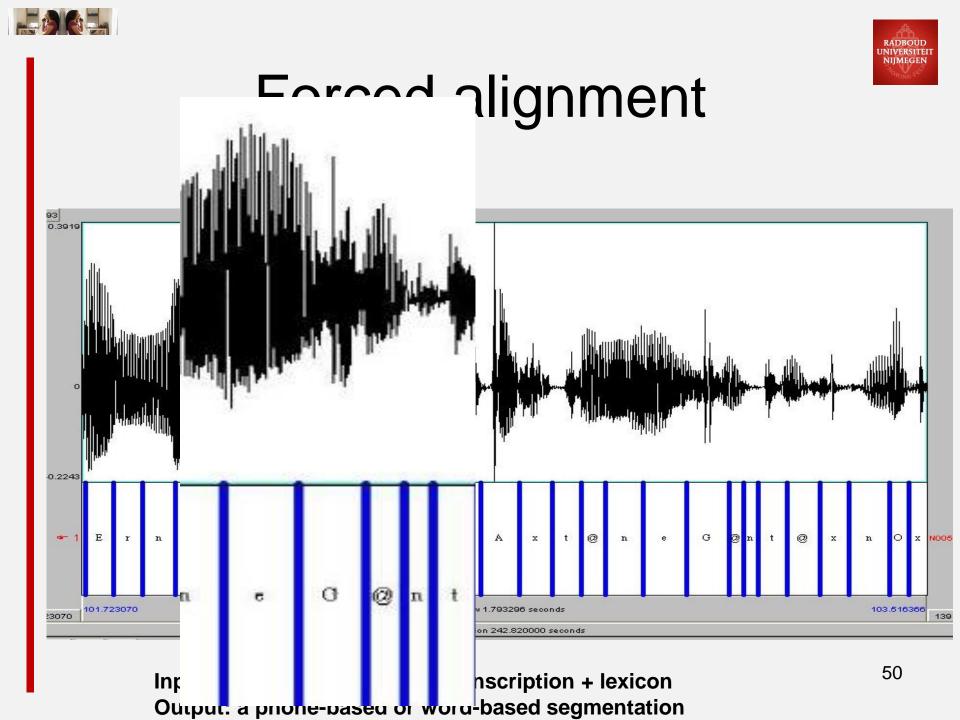




# Forced alignment



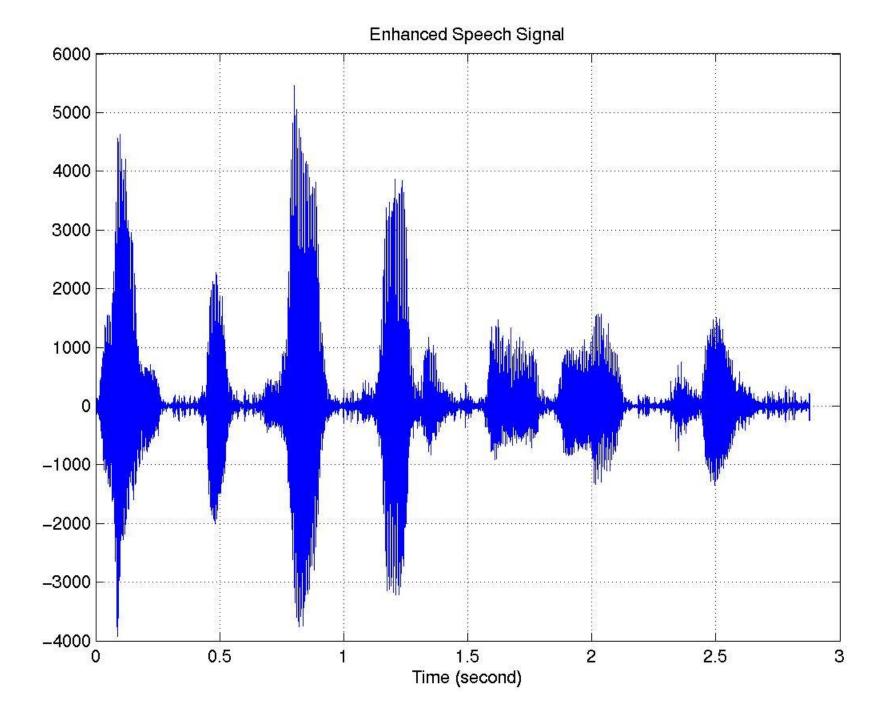
Input: wave file + word-level transcription + lexicon
Output: a phone-based or word-based segmentation

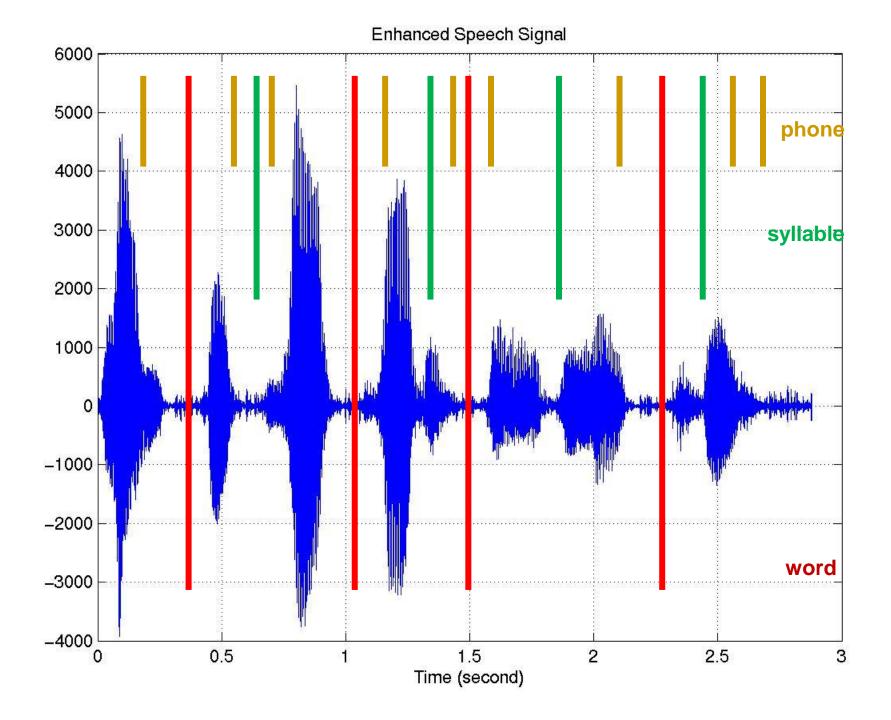




51

Inς nscription + lexicon
Oμιρμι. a prione-pased or word-based segmentation









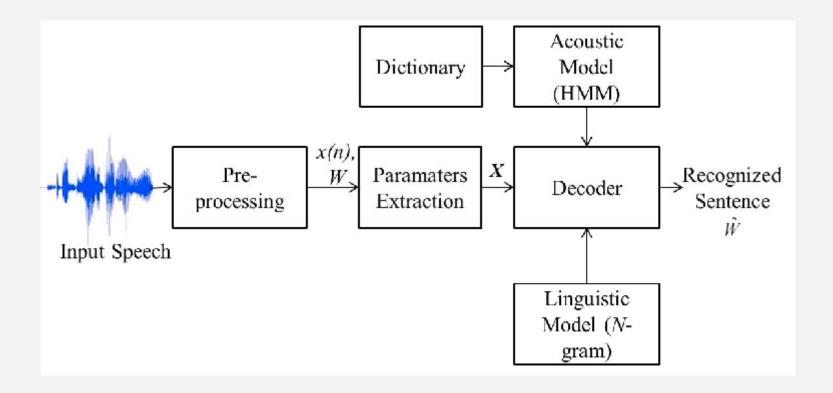
## Small history: ASR

- 1975: recognition using templates
- 1990: Gaussian Mixture Models (GMMs) + Hidden Markov Models (HMMs)
- in 2009 Brno: proto-KALDI system based on FSTs (FST = finite state transducer)
  - https://web.cs.ucdavis.edu/~rogaway/classes/ 120/spring13/eric-transducers.pdf
- since 2012-13: advent of Artificial Neural Networks (ANNs, DNNs, RNNs)





#### ASR: architecture since 1980



Since 2013: DNNs replace or modify one or more components





#### **ASR**

#### Basic principle

HMM *models* of speech units (words, syllables, phones, context-dependent phones)



Signal preprocessing

(front end)

Acoustic

models

Speech recognition

(search algorithm)

Words or phones + alignment

Pronunciation model (lexicon)

Language model





#### **ASR**

#### Basic principle

HMM *models* of speech units (words, syllables, phones, context-dependent phones)

models



Signal preprocessing

(front end)

Acoustic Speech recognition

(search algorithm)

Language model

**Pronunciation** 

model (lexicon)

Words or phones + alignment

N-Best, etc.

3

58





## ASR components

The **front-end**: extracting features from audio

The lexicon specifies per word the corresponding sequence of speech units. There may be pronunciation variants (more variants per word) (word list, FST)

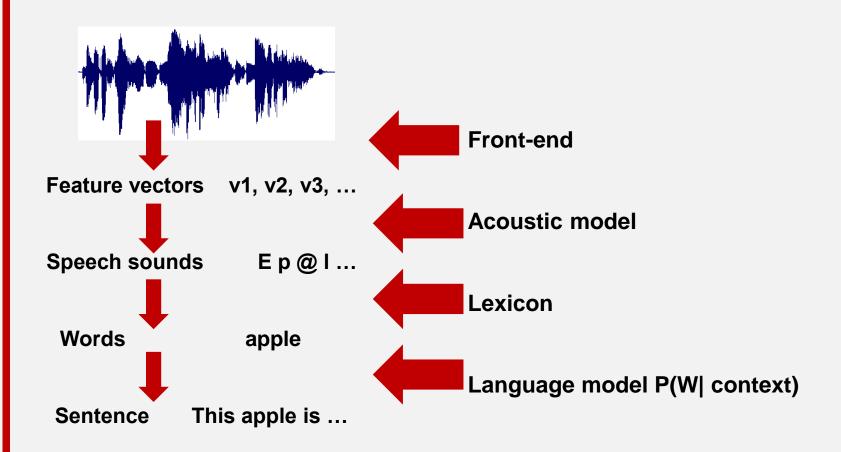
The language model specifies the probability of all possible word sequences. It is derived from large text corpora. The tokens in the LM are the words figuring in de lexicon (N-grams, NN-LM)

The acoustic models describe for each speech unit its internal organisarion (HMM topology) and the probability density function of the corresponding feature vectors (GMM, NN)

The search algorithm decides for the observed frames (under the constraints of acoustic model, lexicon and language model) which word sequence is the most likely one (often FST based)











## Some terminology

Sentence linguistic concept

Utterance a stretch of spoken words

Backchannel uhm, ehm, ja

Word everything between spaces?

Segment linguistic: "segmental tier"

Phoneme smallest unit carrying semantic diff

Phone spoken version of a phoneme

Allophone context dependent phone

Coarticulation influence of neighboring phones



#### Automatic Speech Recognition

- (Usually) based on statistical models
- Given "training data" from the target language, we train a statistical models of speech
  - AM: audio, LM: text
- Given a waveform, we can work out the most likely word sequence via Bayes

P(words | audio) ~ P(audio | words) P(words)



#### Automatic Speech Recognition

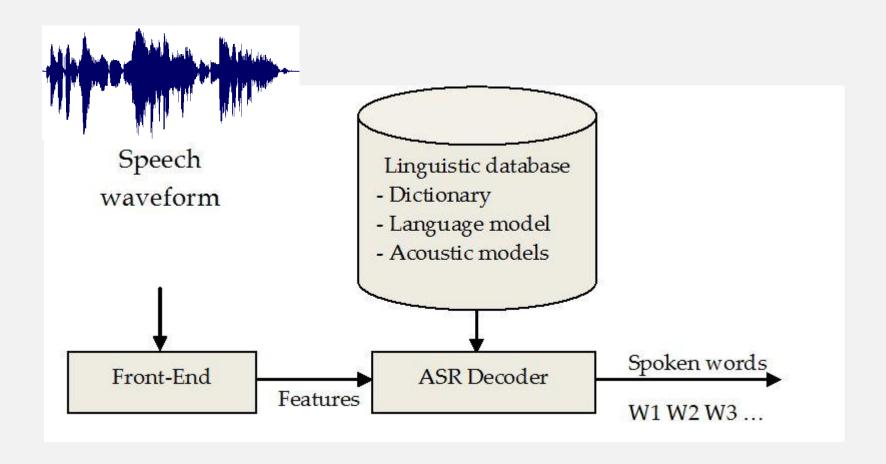
- (Usually) based on statistical models
- Given "training data" from the target language, we train a statistical models of speech
  - AM: audio, LM: text
- Given a waveform, we can work out the most likely word sequence via Bayes







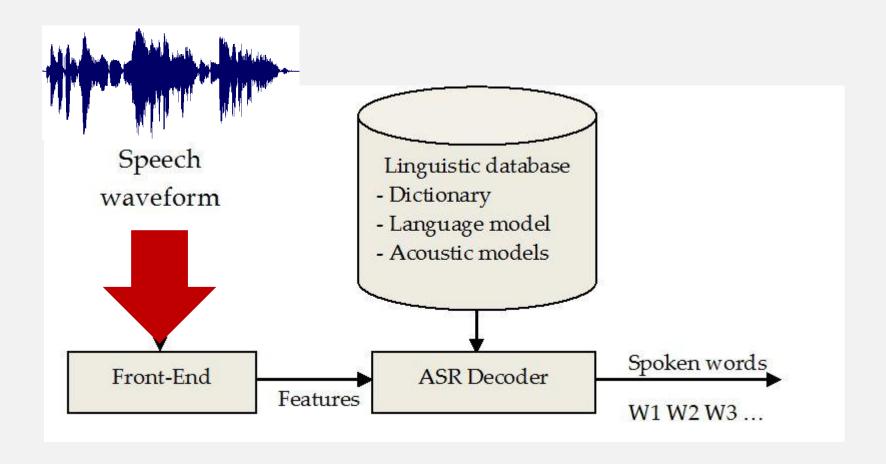
#### ASR architecture







#### ASR architecture



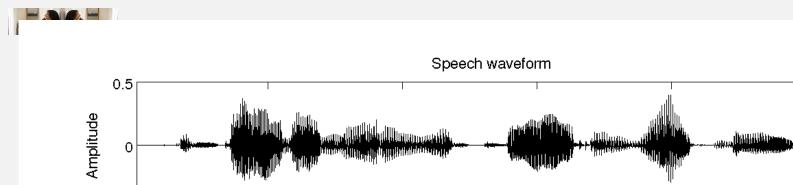




#### From audio to features

- Analog signal to digital signal
- Digital signal to feature vectors
  - -5 steps

- E.g. to MFCC
  - Standard but not unique
  - See "python mfcc"



-0.5 <sup>l</sup>

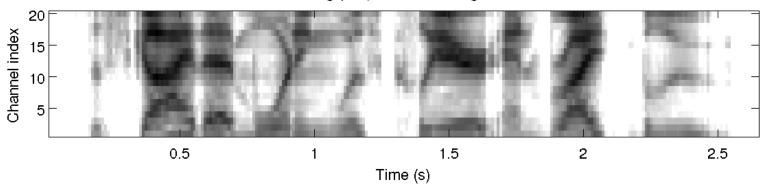
0.5

Time (s) Log (mel) filterbank energies

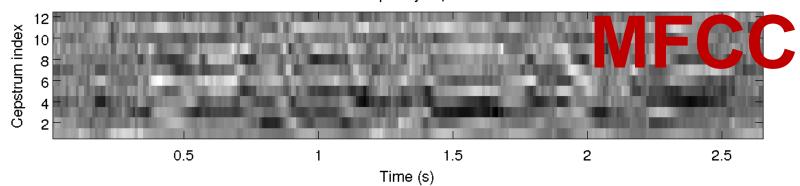
1.5

2

2.5



#### Mel frequency cepstrum

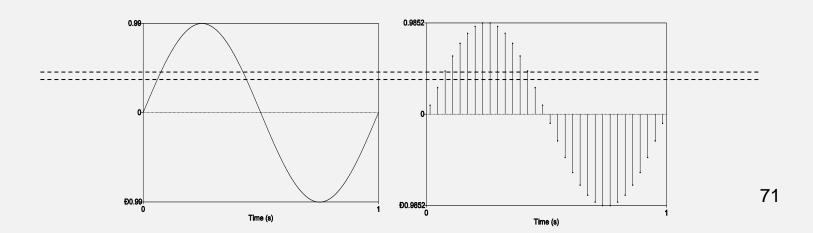






#### Analogue-to-digital (AD) conversion

- Discretisation in time
  - sampling frequency or sampling rate (samples/sec, Hz) determines the highest frequency that can be represented. Nyquist.
- Discretisation in amplitude
  - Number of possible amplitude values is determined by available storage space per sample, e.g.
    - 8 bits (1 byte):  $2^8$  (256) possible values
    - 16 bits (2 bytes): 2<sup>16</sup> (65536) possible values
  - Amplitude values are rounded to nearest discrete value







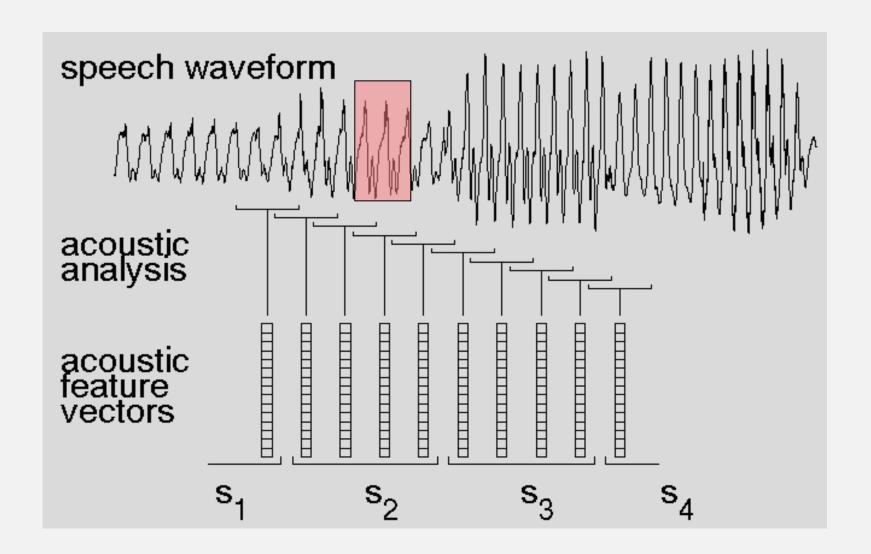
# Signal to features

# 5 steps to convert a digital speech signal into acoustic feature vectors:

- 1. Segmentation: Division of the speech signal (oscillogram) into segments with a sufficiently short duration to assume that the vocal tract does not change shape (*frame-based analysis*)
- 2. Apply windowing
- 3. Convert the windowed signal from the time domain to the frequency domain using the *Fast Fourier Transform (FFT)*
- 4. Apply perceptual weighting on the spectral energy distribution (e.g. Mel or Bark scale) so that the energy in a certain frequency band becomes approximately equally important for the computer as for a human listener.
- 5. Decorrelate the output of the filters so that the histograms of the resulting features can be modeled adequately by mutually independent Gaussian probability density functions.



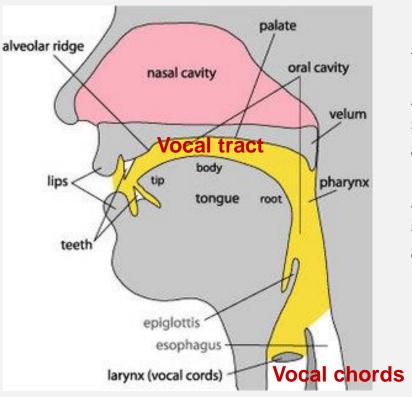








### Segmentation: vocal tract



**Articulation is relatively slow** 

About 12-14 speech sounds per second, i.e. 70 ms. per phone, on average

Articulations move synchronously/in parallel → assimilation of properties of neighboring sounds



# Segmentation: analysis window

A long analysis window (e.g. 50 ms)

- → very precise description of how energy is distributed over frequency (good spectral resolution)
- → but: poor temporal resolution
- → good but only useful if speech sounds last long enough: e.g. vowels, fricatives

A short analysis window (e.g. 10 ms) allows a detailed tracking of fast changes → good time resolution (useful for an accurate description of plosives) but poor spectral resolution

In practice, window length is often 25 ms (2.5 glottal periods for males; about 5 glottal periods for females)





# Segmentation: duration

Segment the speech signal in small intervals (frames)

What is being said is mainly determined by the shape of the vocal tract. The shape of the vocal tract determines the energy envelope spectrum. <a href="https://www.youtube.com/watch?v=cVovOKLISb0">https://www.youtube.com/watch?v=cVovOKLISb0</a>

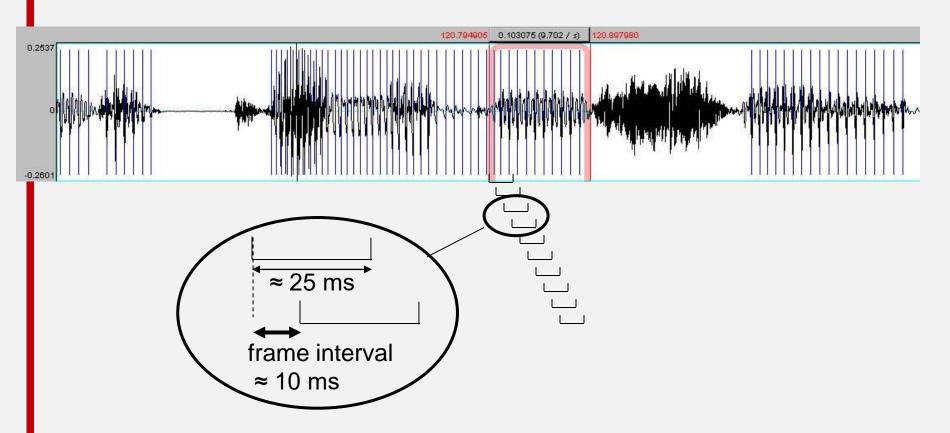
What is a sensible analysis duration for a frame? The time interval in which the vocal tract does not change substantially is determined by the number of speech sounds per second that is produced (≈13-15). (70ms) A defendable analysis frame duration is 25 ms.

How many frames per second? To accurately describe the changes in vocal tract shape over time, the number of frames per second must be at least twice as high as the highest frequency with which the vocal tract changes (Nyquist criterion) → 100 times/second (i.e. every 10 ms).





### Feature extraction



25ms analysis window length 10ms frame shift

If sample freq = 16kHz,  $25ms \sim 0.025*16000 = 400 samples$ .





### Feature extraction

# 5 steps to convert a digital speech signal into acoustic feature vectors:

- 1. Divide the speech signal (oscillogram) in segments with a sufficiently short duration to assume that the vocal tract does not change shape (frame-based analysis)
- 2. Apply windowing
- 3. Convert the windowed signal from the time domain to the frequency domain using the Fast Fourier Transform (FFT)
- 4. Apply perceptual weighting on the spectral energy distribution (e.g. Mel or Bark scale) so that the energy in a certain frequency band becomes approximately equally important for the computer as for a human listener.
- 5. Decorrelate the output of the filters so that the histograms of the resulting features can be modeled adequately by mutually independent Gaussian probability density functions.





# Feature extraction: windowing

Taking a "rectangular" unweighted portion out of a wave file leads to audible artefacts. These artefacts can be avoided by proper windowing: taper off the beginning and end of the signal.

For a well-chosen window, the spectrum is nearly identical to a signal of which the core part is repeated indefinitely.

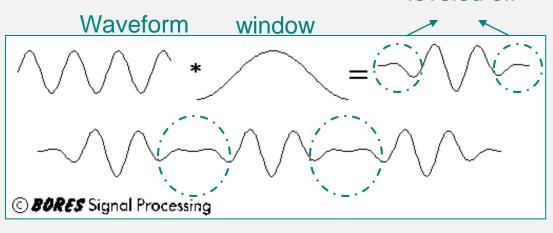
Often used windows are Hamming and Hanning windows. See e.g. <a href="https://www.youtube.com/watch?v=YsqGQzJ\_2V0">https://www.youtube.com/watch?v=YsqGQzJ\_2V0</a>

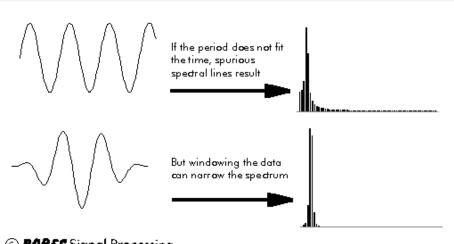




# Windowing

#### leveled off









#### Feature extraction

# 5 steps to convert a digital speech signal into acoustic feature vectors:

- 1. Divide the speech signal (oscillogram) in segments with a sufficiently short duration to assume that the vocal tract does not change shape (frame-based analysis)
- 2. Apply windowing
- 3. Convert the windowed signal from the time domain to the frequency domain using the Fast Fourier Transform (FFT)
- 4. Apply perceptual weighting on the spectral energy distribution (e.g. Mel or Bark scale) so that the energy in a certain frequency band becomes approximately equally important for the computer as for a human listener.
- 5. Decorrelate the output of the filters so that the histograms of the resulting features can be modeled adequately by mutually independent Gaussian probability density functions.

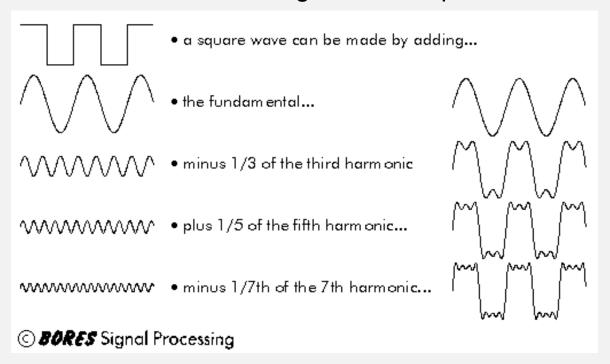




#### Feature extraction: FFT

Fast Fourier transform (FFT): from time domain to frequency domain

 Jean-Baptiste Fourier: Every waveform is the sum of sine waves with a certain magnitude and phase

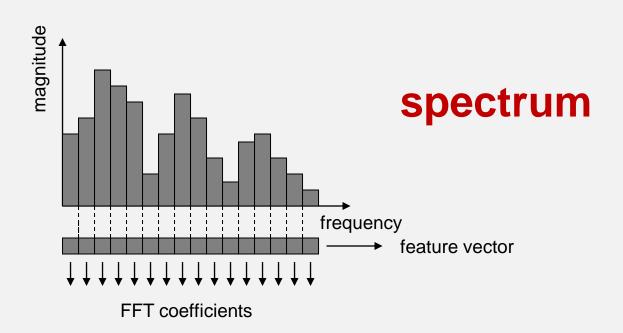






#### FFT

A Fast Fourier Transform (FFT) decomposes the speech signal within each window as a weighted sum of complex exponentials









The resulting FFT coefficients are written in one *vector* (one per frame).

The more coefficients, the more accurate the description, but higher order coefficients may be noisy.





#### Feature extraction

# 5 steps to convert a digital speech signal into acoustic feature vectors:

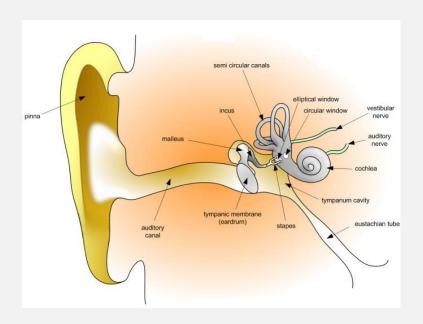
- 1. Divide the speech signal (oscillogram) in segments with a sufficiently short duration to assume that the vocal tract does not change shape (frame-based analysis)
- 2. Apply windowing
- 3. Convert the windowed signal from the time domain to the frequency domain using the Fast Fourier Transform (FFT)
- 4. Apply perceptual weighting on the spectral energy distribution (e.g. Mel or Bark scale) so that the energy in a certain frequency band becomes approximately equally important for the computer as for a human listener.
- 5. Decorrelate the output of the filters so that the histograms of the resulting features can be modeled adequately by mutually independent Gaussian probability density functions.



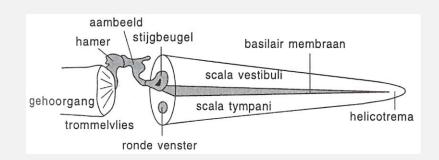
# Feature extraction: Perceptual weighting



Convert FFT coefficients into a set of features using a perceptual weighting inspired by the processing characteristics of the human auditory system.



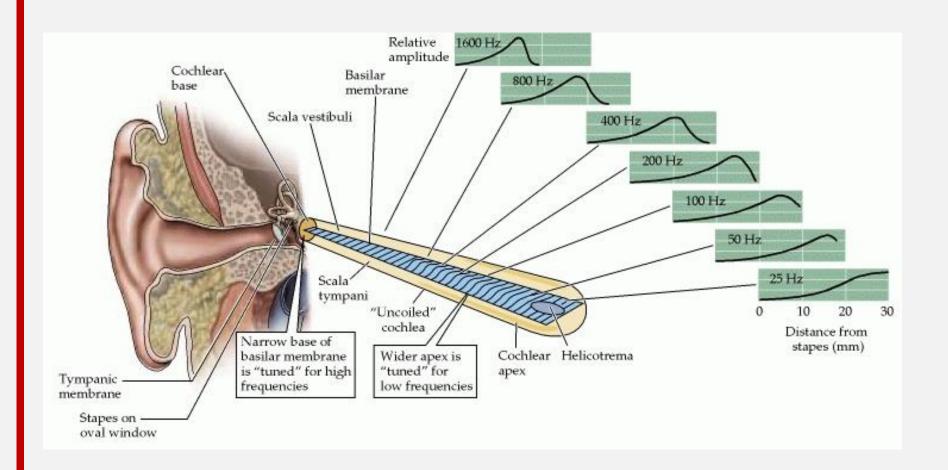
#### malleus, incus, stapes



Cochlea in normal (left, wikipedia) en unrolled form











#### Cochlea

- The cochlea: tube filled with fluid of which the diameter gradually tapers off
  - It is divided into 2 compartiments by the basilar membrane which decreases in thickness
- Approx. 30,000 hair cells attached to the basilar membrane
- A sound makes the basilar membrane vibrate. For a pure tone (a sinusoide) the oscillations are largest at one specific location of the basilar membrane.
- This peak location is determined by the frequency of the tone. At the end: low freqs. At the beginning: high freqs.
  - → TONOTOPIC mapping





# Perceptual weightings

- The human ear is not sensitive to frequency along a linear scale
- Psycho-acoustical frequency scales often used to approximate the human sensitivity are the *Mel scale* and the *Bark scale*

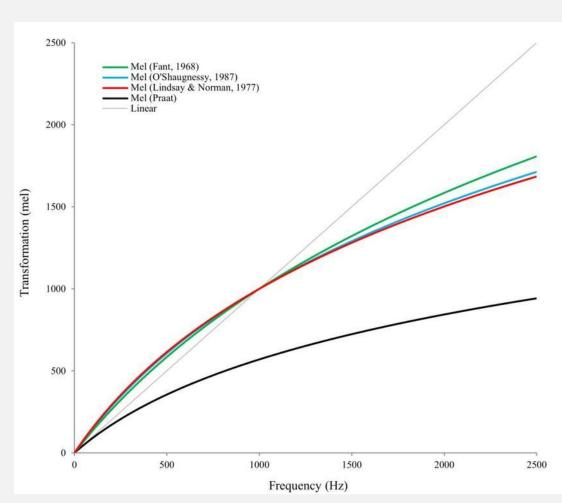




# frequency to mel

mel(f) = 1125 log(1+f/700)

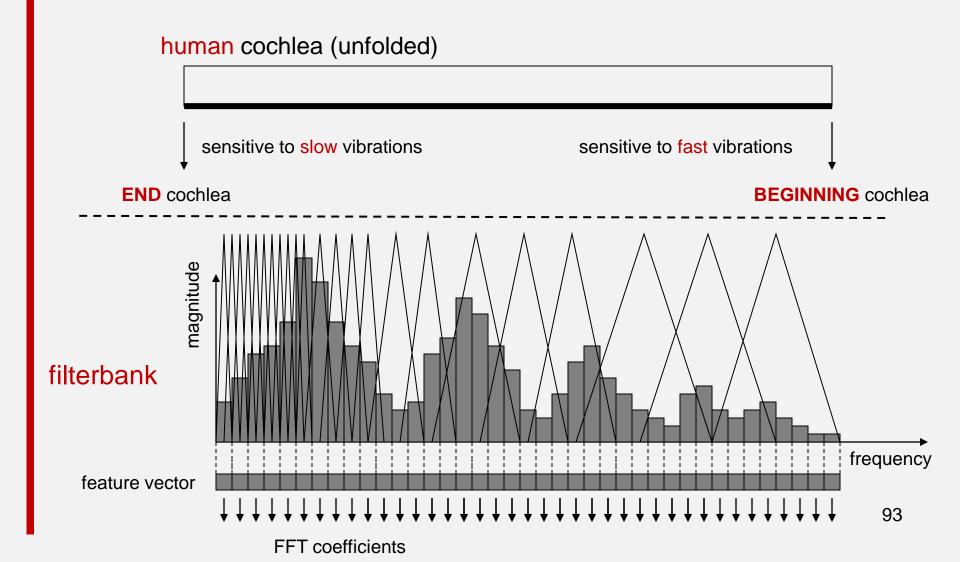
There are several other transformations, all log-like.







### Mel-filterbank





# Apply log() to all energy values in each filter



•  $E \rightarrow log(E)$ 

$$\Delta Percept = \frac{\Delta Physical Quantity}{Physical Quantity}$$

- Weber's law
- https://www.youtube.com/watch?v=hHG8io5qIU8
  - Energy → loudness
  - Fundamental frequency → pitch (piano)
  - Perception of physical phenomena
  - Estimation of physical quantities
  - Duration, length, pressure, ...





### Feature extraction

# 5 steps to convert a digital speech signal into acoustic feature vectors:

- 1. Divide the speech signal (oscillogram) in segments with a sufficiently short duration to assume that the vocal tract does not change shape (frame-based analysis)
- 2. Apply windowing
- 3. Convert the windowed signal from the time domain to the frequency domain using the *Fast Fourier Transform (FFT)*
- 4. Apply perceptual weighting on the spectral energy distribution (e.g. Mel or Bark scale) so that the energy in a certain frequency band becomes approximately equally important for the computer as for a human listener.
- Decorrelate the output of the filters so that the statitics of the resulting features can be modeled by mutually independent Gaussian probability density functions.





#### Decorrelation

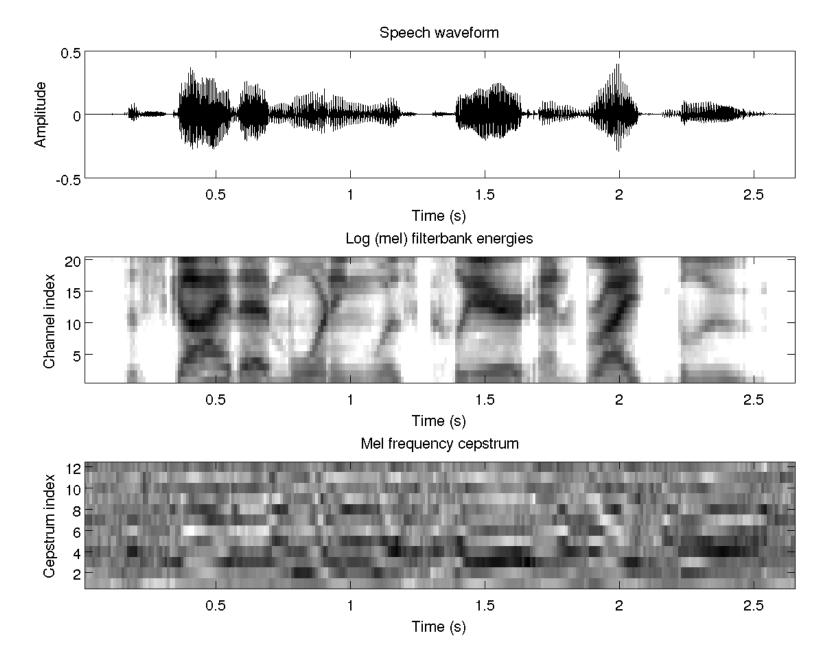
The amount of energy in neighbouring filters is strongly correlated. In order to reduce this correlation, a Discrete Cosine Transform (DCT) is performed on the 21 filter bank coefficients.

As a result we obtain the **Mel Frequency Cepstral Coefficients (MFCCs).** 

These MFCCs are approximately statistically independent. The first 12 coefficients  $c_1...c_{12}$  suffice to describe the relevant details of the spectrum within each frame. See e.g.

https://haythamfayek.com/2016/04/21/speech-processingfor-machine-learning.html for comments on this DCT step









# Signal preprocessing

#### **Summary**

0. A/D co	nversion
-----------	----------

- 1. sliding portion
- 2. windowing
- 3. FFT
- 4. filterbank
- 5. MFCC

	(
	typical ciza
output	typical size
	-

digital signal -

analysis stretch 400 samples

windowed signal 400 samples

spectrum 400 magnitudes

feature vector 21 filterbank energies

smaller feature vector 12 features

assuming 16kHz, 25 ms analysis frame



# MFCC vectors as audio representation



Very many sites show information about MFCCs, often with useful Python function calls

See e.g.

https://www.kaggle.com/ilyamich/mfccimplementation-and-tutorial

https://pypi.org/project/python\_speech\_features/

Librosa python library

https://librosa.org/doc/latest/index.html

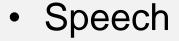


# Difference between image and speech recognition/classification



#### Images

- Static (often)
- Size 400 x 400 (or so)
- Foreground hides background
- Focus



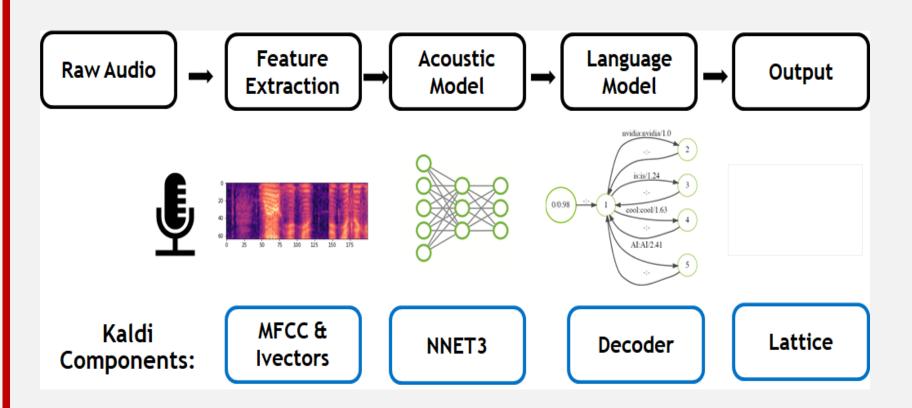
- No clear segmentation
- Unfolding over time (100 spectral frames/s)
- Words are not immediately available
- All phones are context/speaker/mood/ ... dependent







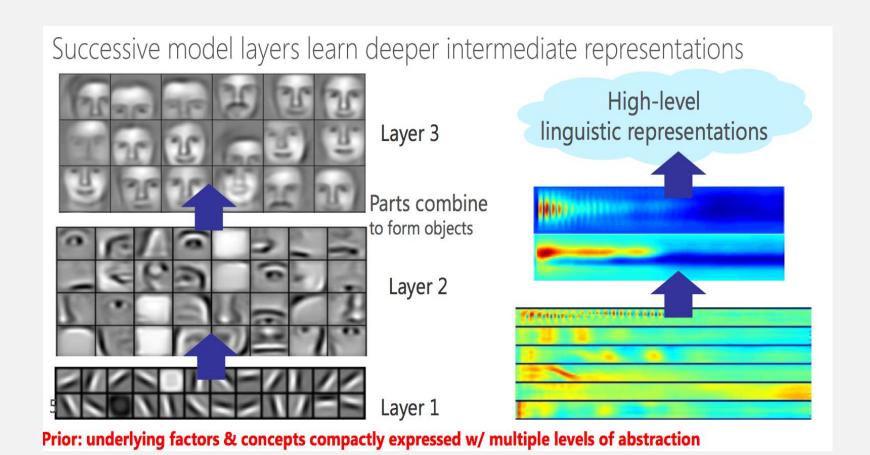
#### Classical and new (deep) learning



https://devblogs.nvidia.com/gpu-accelerated-speech-to-text-with-kaldi-a-tutorialon-getting-started/



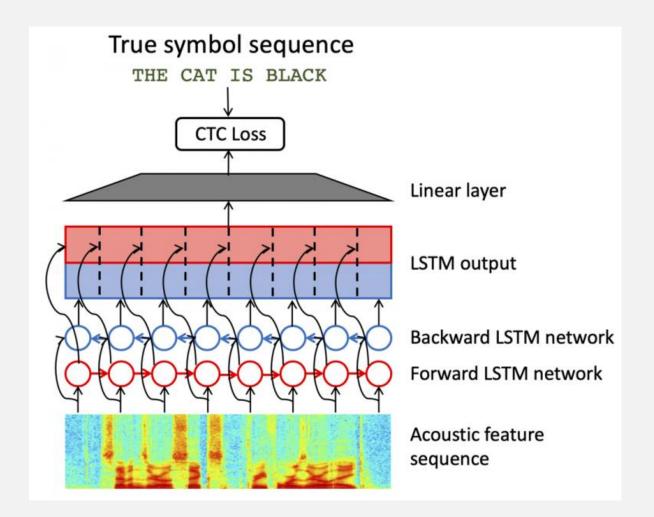




Lee, Largman, Pham & Ng, NIPS 2009







Nowadays, ASR is an example of deep learning





	pro	con
Classical approach	Insight-based	Lower performance
Recent deep learning approaches	Higher performance (relative 30-60% reduction error rates)	What do <b>we</b> learn? Explainable Al Open Al Responsible Al





### Deep learning in ASR

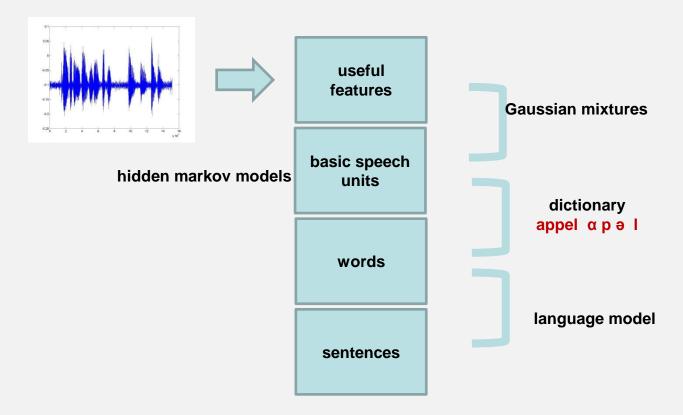
 It is very useful to know about the conventional approach before diving into DNN-based approaches

 Nowadays hundreds of (highly specialized) papers on DNNs in ASR





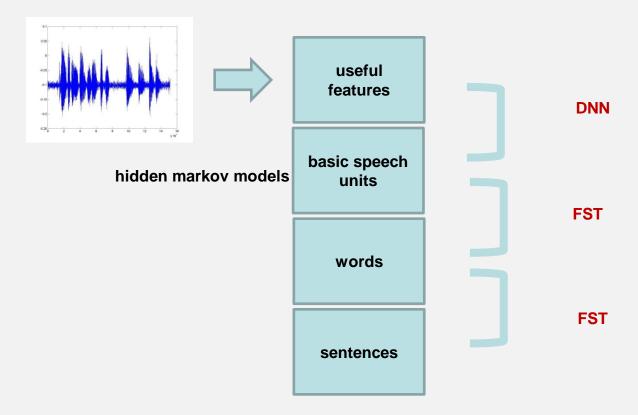
#### conventional ASR







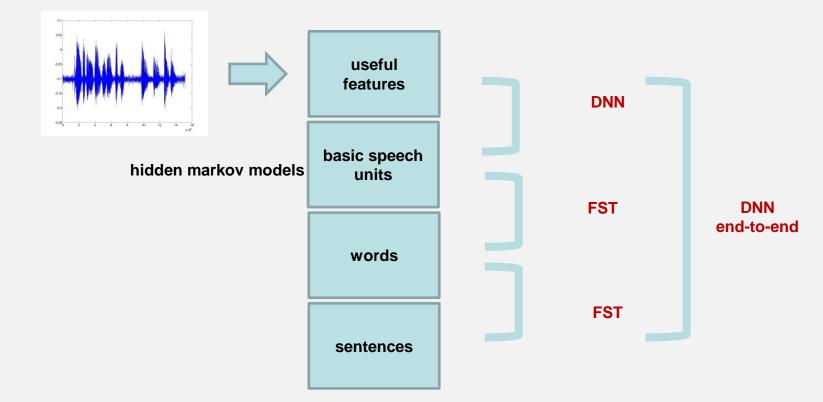
#### modern ASR







#### modern ASR



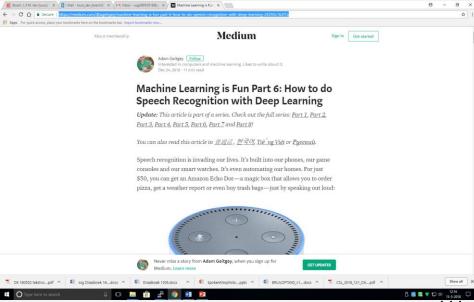




# ASR and deep learning

 https://medium.com/@ageitgey/machinelearning-is-fun-part-6-how-to-do-speechrecognition-with-deep-learning-

28293c162f7a



111





### ASR techniques

#### Template matching ('70)

- Representative speech patterns (e.g. words, syllables) are stored
- And labeled
- New unknown data identified by comparison

#### Probabilistic matching ('80-now)

- Often by Hidden Markov Modeling (HMM)
- An HMM is characterized by hidden states graph
   + state-state transition probabilities
- Each state is characterized by a (statistical) distribution in the acoustic space

#### ANN-based (2012-now)

Many directions, most of them based on earlier approaches





## Template matching

- Used 1960-70 for isolated words
  - Record a database of words from same speaker (or different speakers)
  - In the test, compare a test item to all of the stored speech tokens
- Distance: Dynamic Time Warping
  - DTW: a way to define a "distance" between sequences of unequal length (see later)

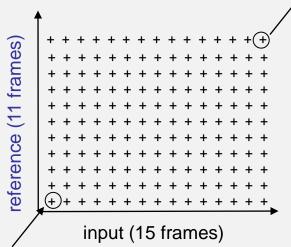




# Template based recognition

- By dynamic programming
- Define a trellis matrix (or lattice)
- First pass: in each node (cell), compute all local matching scores, compute forward global scores and keep pointers

Second pass: backtrace

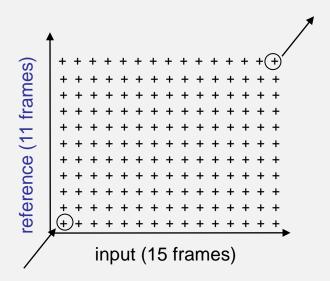


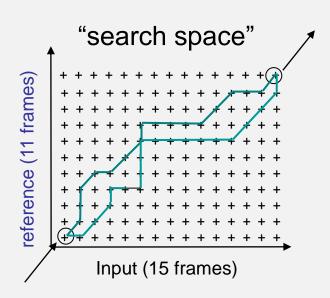






This algorithm provides the best alignment between input and reference sequences







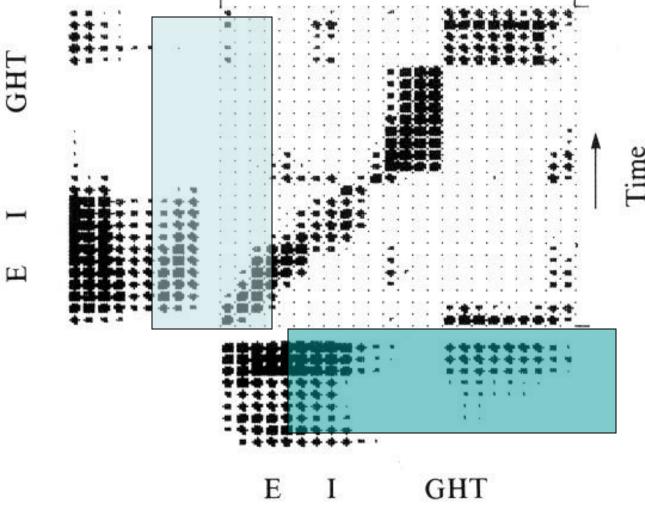


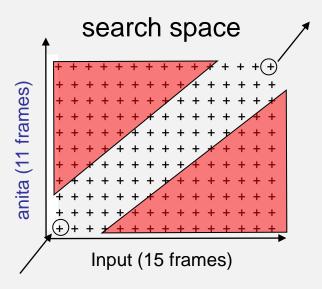
Fig. 7.2 Graphical representation of Euclidean distance between frames of the spectrograms shown in Fig. 7.1. The larger the blob the smaller the distance. It can be seen that there is a path of fairly small distances between the bottom left and top right when two examples of 'eight' are compared,





#### pruning

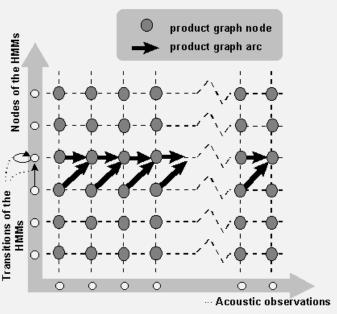
Computational costs can be further reduced by pruning: prohibit too unlikely alignments by applying local scores that prohibit/penalize implausible search areas

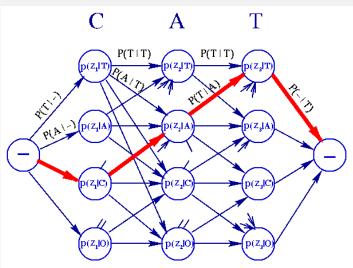


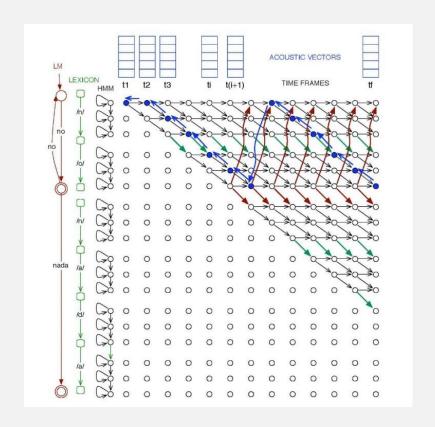




#### Other examples of dynamic search







## 1980-1990: towards probabilistic models



Three main factors played a role in ASR advancement

- more speech data & more disk space
- More powerful CPU/GPU
- better algorithms for estimating Hidden Markov Models parameters (Baum-Welch, 1972)

Exactly these factors also play a role in the current network approaches in ASR.

# Bayes plays key role in almost all approaches

$$P(A \mid B) = \frac{P(A \cap B)}{P(B)}$$
 Definition of P(A|B) 
$$P(A \cap B) = P(A \mid B)P(B)$$
 
$$P(A \mid B)P(B) = P(B \mid A)P(A)$$
 
$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}$$

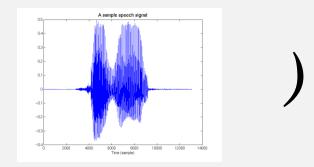
# Bayes plays key role in almost all approaches

$$P(A \mid B) = \frac{P(A \cap B)}{P(B)}$$
 Definition of P(A|B) 
$$P(A \cap B) = P(A \mid B)P(B)$$
 
$$P(A \mid B)P(B) = P(B \mid A)P(A)$$
 
$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}$$

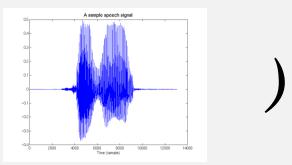




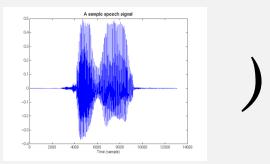




P(I sad/



P(I sat/







```
P(I said | utterance)
```

P(I sad | utterance)

 $P(word(s) \mid utterance)$ 





```
P(I said | utterance)

P(I sad | utterance)
```

P(woods) utterance)





## P(word(s) | utterance)

$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}$$

$$P(word \mid utterance) = \frac{P(utterance \mid word)P(word)}{P(utterance)}$$





## Bayes

$$P(word \mid utterance) = \frac{P(utterance \mid word)P(word)}{P(utterance)}$$

 $arg \max P(word_i | utterance) =$   $arg \max P(utterance | word_i) P(word_i)$ 

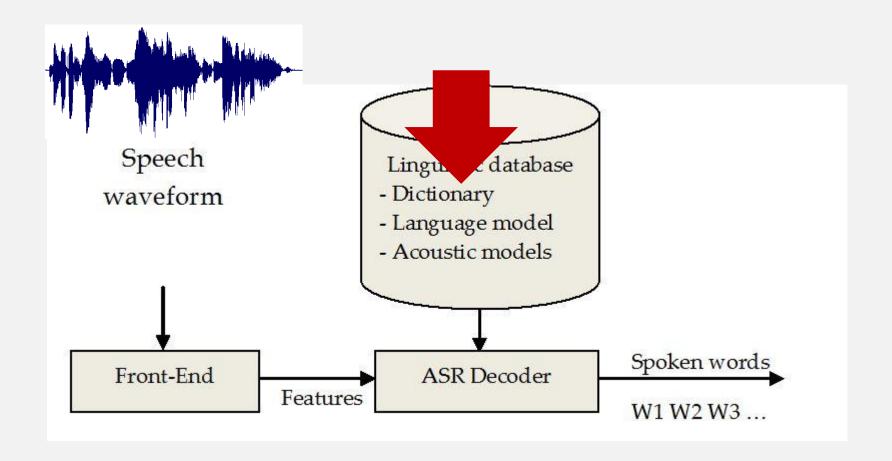
**AM** 

LM





#### ASR architecture







## Lexicon/dictionary

word form	pł	nor	<u>ien</u>	<u>nic</u>	tra	ns	<u>cri</u> p	otic	<u>n</u>
bang	b	A	N						
bankbiljet	b	A	N	b	I	1	j	E	t
barbaren	b	A	r	b	a	r	9		
barbecue	b	A	r	b	9	k	j	u	W
onmiddellijk	0	m	I	d	<u>@</u>	1	9	k	
yesterday	j	E	s	t	9	r	d	е	
goedemorgen	x	u	9	m	0	r	X		





## Lexicon/dictionary (Dutch)

word formphonemic transcriptionbangb A Nbankbiljetb A N b I l j E tbarbarenb A r b a r @barbecueb A r b @ k j u wonmiddellijkO m I d @ l @ k

For a specific language, lexicons are often available, but not always a clean source of information.

j E s t @ r d e

Multiples used?

yesterday

Phonetic alphabet used?

Compounding? (Rechtsschutzversicherungsgesellschaften)





#### Lexicons...

A pronunciation lexicon lists for every word the sequence of phonetic symbols that describes its pronunciation

- Pronunciation mostly in terms of broad phonetic transcription
- Usually no lexical stress
- Usually no additional info, such as POS, morphological properties

#### Two types:

- Canonical lexicon: one pronunciation/word (norm pronunciation)
- Multiple pronunciation lexicon: different pronunciations/word to account for possible pronunciation variation





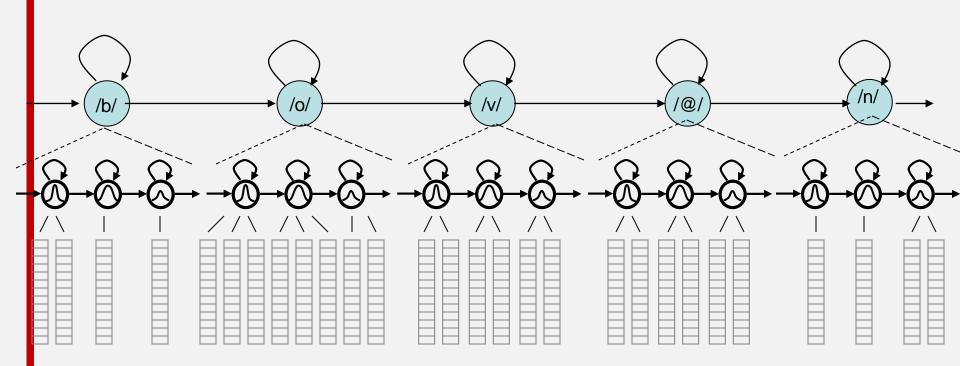
	Pro	Con
Canonical	simple	actual pronunciation may mismatch (reduction phenomena)
Multiple pronunciation	better description of actual realization of the speech signal	Increased confusion across words (homophones)





## Modeling words from phones

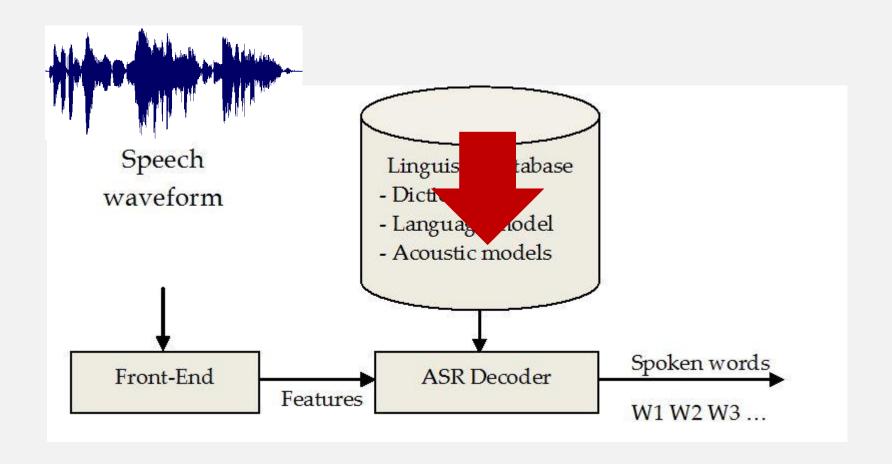
A word can be represented as a concatenation of phone models







#### ASR architecture





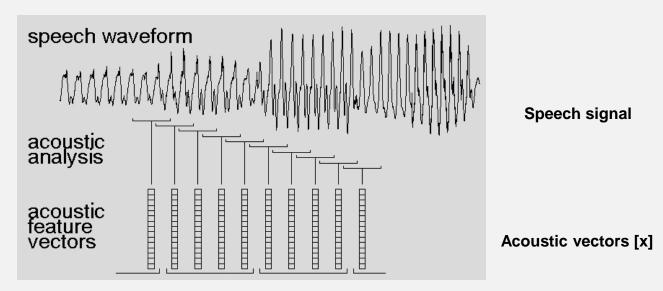


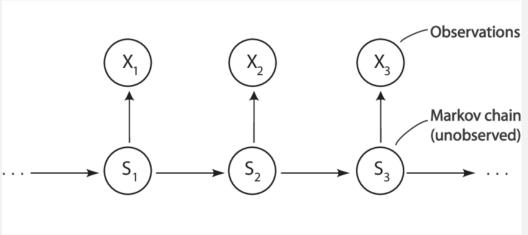
#### Acoustic model

- The AM describes the statistics of speech sounds, for each speech unit
- Different speech units can be used, e.g.:
  - Phones (monophones): p, t, k, f, v, a, i, u,...
  - Triphones, quintphones, ...
  - Syllables
  - Words
  - Sentences
  - Etc.





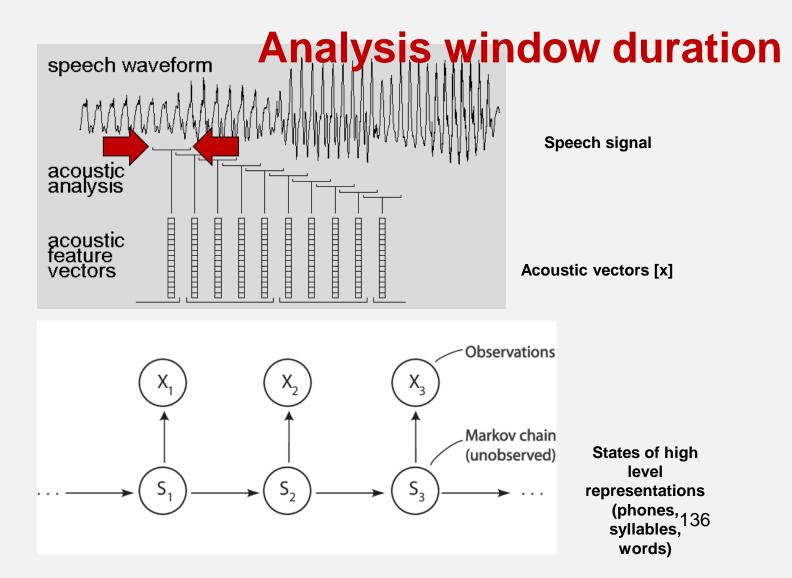




States of high level representations (phones, 135 syllables, words)

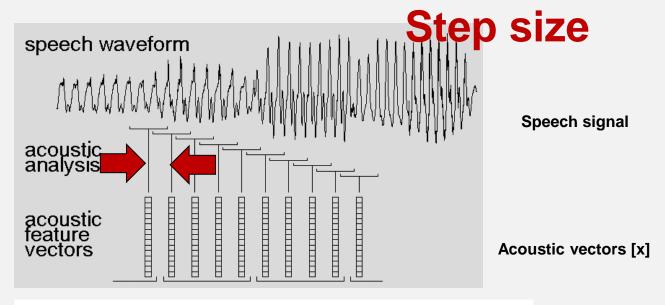


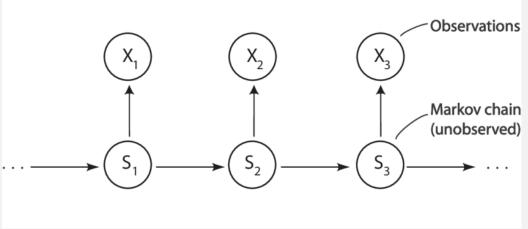












States of high level representations (phones, 137 syllables, words)





#### Acoustic models

- Describe the statistics of the corresponding acoustic feature vectors per speech unit
- Different speech units can be used:
  - Sentences
  - Words
  - Syllables (e.g., Mandarin)
  - Phones
- Mostly implemented as HMMs (3-state)
- The most frequently used units are phones (because very generic), triphones, quintphones
  - Triphone = 1-left-1-right context-dependent phone
  - Quintphone = 2-left-2-right context-dependent phon





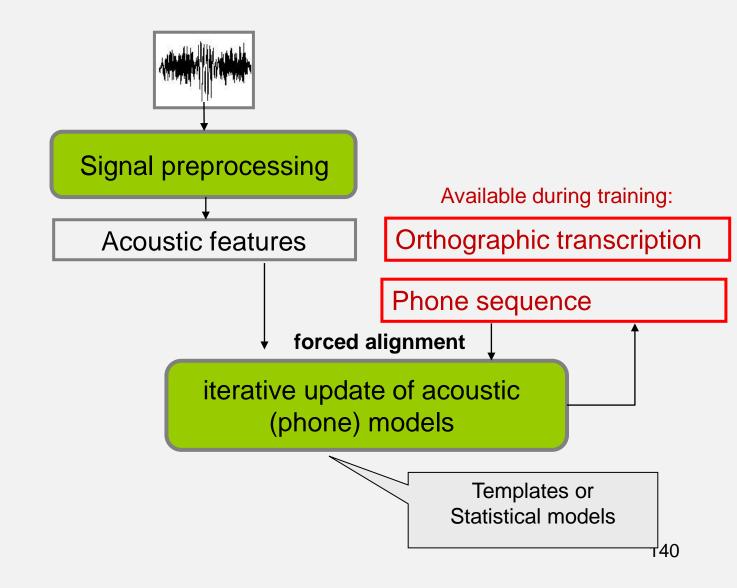
#### Some numbers

- Languages use about 30 50 phones
- An AM contains 30 50 context independent phones
  - 'apple' E p @ I
- May contain 2000 3000 triphones
  - 'apple'  $_{\#}E_{p} _{E}p_{@p}@_{l} _{@}I_{\#}$
- Several other options





## ASR in a diagram: training







## Model based ASR: training

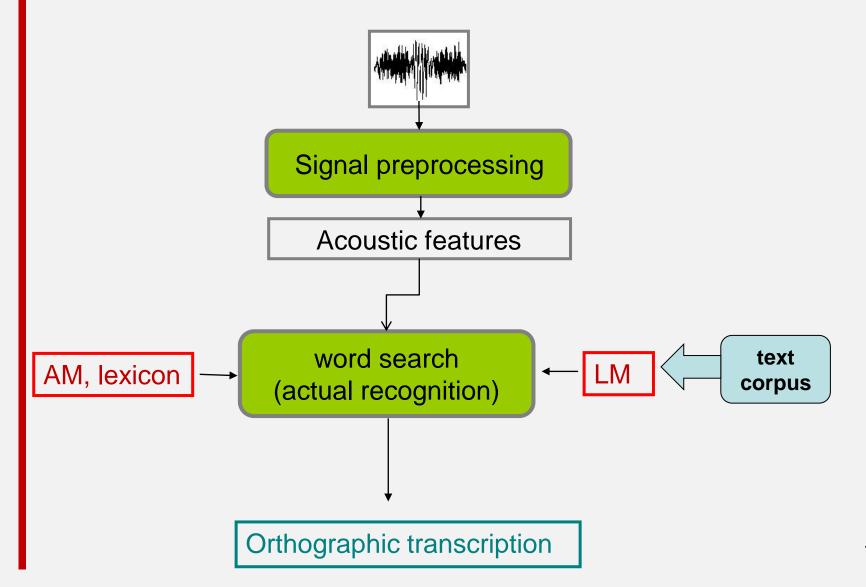
#### (simplified)

- 1. Starting material: A training corpus with:
  - Speech
  - Orthographic transcriptions
  - Pronunciation dictionary (lexicon)
- 2. Create lousy acoustic HMM (phone) models
- 3. The transcriptions are "forced aligned" with the speech signal (no LM)
- 4. Corresponding HMM models are updated
- 5. Back to 3, until convergence





## ASR in a diagram: test



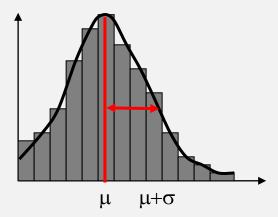




## Statistics per HMM state

Via DNN, or via a Gaussian mixture

For every state of the HMM, the distribution of the corresponding feature vectors is obtained by using the frame-to-state relations from a forced alignment using previous HMMs



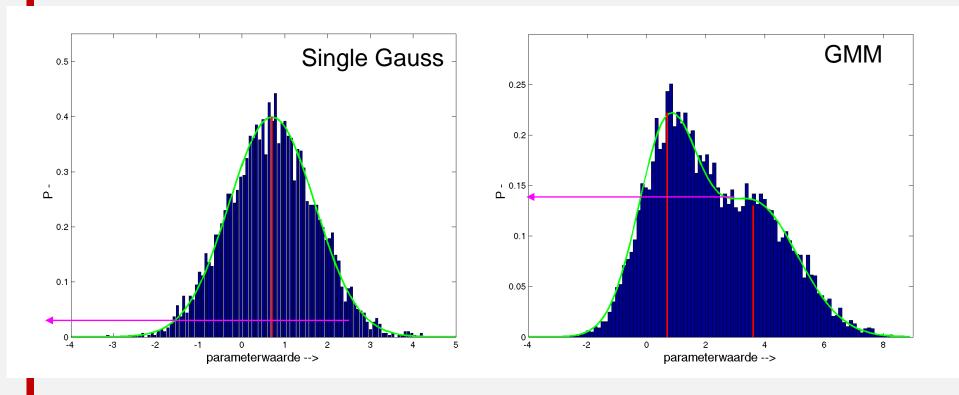
Assuming a normal distribution

Mean (μ) and stdev (σ)





#### Model based ASR



Using a single Gaussian distribution is usually not adequate for modeling the frame distribution.

Instead one may use a weighted sum of Gaussians (a Gaussian mixture, "GMM"), or a DNN.



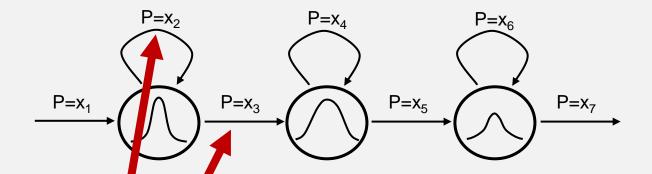


- Hidden: The sequence of states which caused the observed sequence of feature vectors is not an observable.
- Markov property: the state of a system at time t only depends on the state the system was in at time t-1.
- Model: The acoustic properties of a stationary speech segment are described with an acoustic model. Its parameters are emission and transition probabilities.





A typical phone model consists of three states

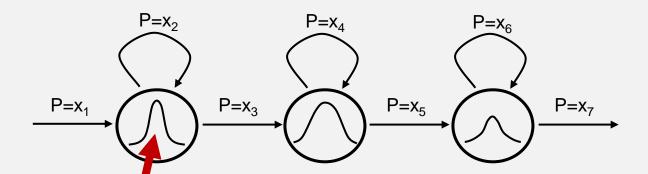


- States Each state models a quasi-stationary part of the spectrogram. It is assumed to generate a feature vector every 10 ms which is drawn randomly according to a state-dependent distribution
- Transitions from one state to the same or the next (or another future state) Transition probabilities are learned during training





A typical phone model consists of three states

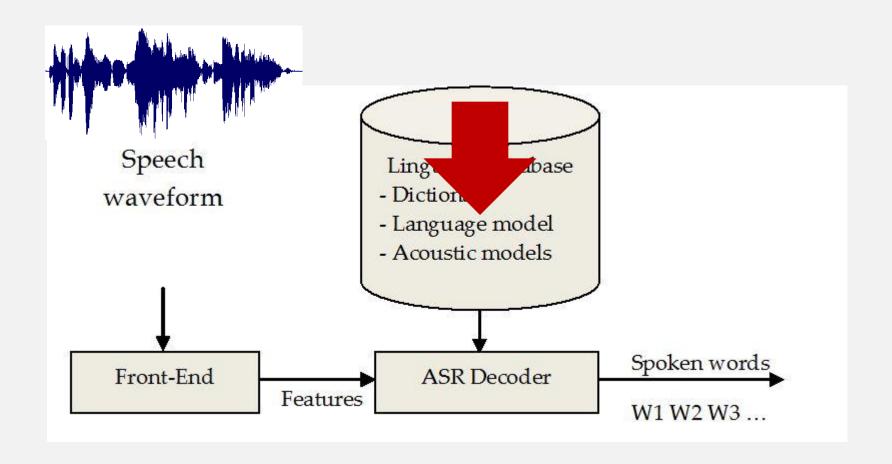


- States Each state models a quasi-stationary part of the spectrogram. It is assumed to generate a feature vector every 10 ms which is drawn randomly according to a state-dependent distribution
- Transitions from one state to the same or the next (or another future state) Transition probabilities are learned during training
- Emission probabilities





#### ASR architecture







## Language model (LM)

LM = A statistical description of "all" sentences

```
P("let's go to the party") =
P("<BOS> let's go to the party <EOS>") =
P("let's | "<BOS>") *
P("go | "<BOS> let's") *
P("to" | ""<BOS> let's go") *
P("the" | "<BOS> let's go to") *
P("party" | "<BOS> let's go to the") *
P("<EOS>" | "<BOS> let's go to the party")
```



**Thomas Bayes** 





#### Statistical LM in ARPA format

-10log(p) n-gram [back-off]

ARPA = Advanced Research
Agency Projects

```
\data\
ngram 1=37445
ngram 2=138797
ngram 3=51837
ngram 4=53201
\1-grams:
-1.727079 A
              -1.184703
-5.57354 AACHEN
                   -0.30103
-5.57354 AAMI -0.30103
-4.833177 AARON
                   -0.3245111
-5.27251 AARONS -0.39794
-5.57354 AARRON
                   -0.30103
-0.4887864
              INCLUDE A -0.8016323
-1.698039
              INCLUDE ART
-1.698428
              INCLUDE BODY
-0.1107919
              ALSO INCLUDE A -0.4771213
-0.1641603
              TO INCLUDE A
                             -1.079181
-0.611348
              TO INCLUDE EVERY
                                  -0.69897
-0.153755
              ALSO INCLUDE A TEN
-0.01548983
              TO INCLUDE A NEW
-0.102658
              WILL INCLUDE A SHIFT
-0.01771067
              IT INCLUDES A STOPOVER
\end\
```





## Language models

 Unweighted language models: a rule based formalism to describe which sequences of words form legal sentences:

we [want to] go home [tonight]

 Probabilistic language models: a formalism that estimates probabilities of word sequences from text:

P(word | precontext)

P("tonight" | "to go home")

Can be a very long list of N-grams

Vectorized language models: e.g., word2vec, in combination with ANNs

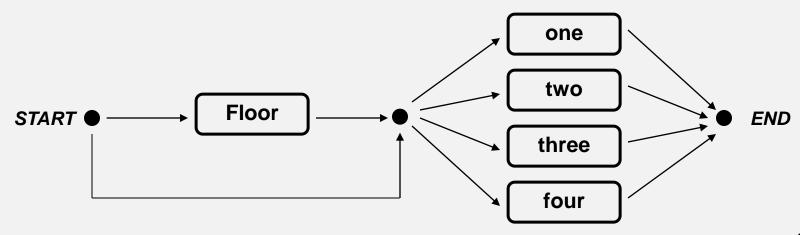
each word is represented as a vector (dim 200-300)





## Language models

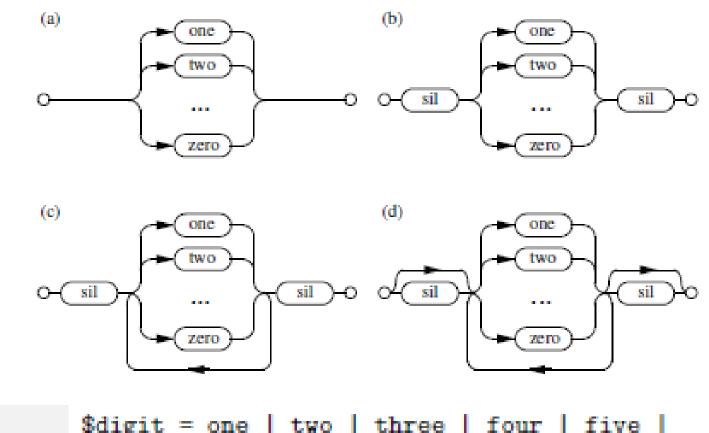
- Unweighted language models:
  - used for applications using low-end memory
  - often a regular (finite state) grammar (→ regexp)
  - can be constructed by hand for a limited domain
  - Example:







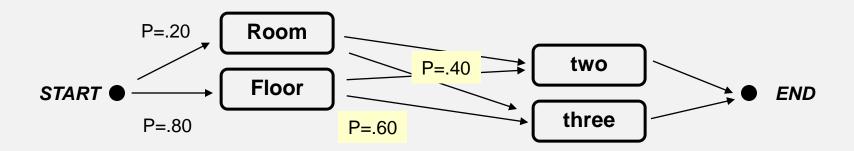
## Other regular grammars







- Mostly used in tasks where more memory is available.
- Estimated by means of orthografic transcriptions from a text corpus







- models P(current word | precontext)
- in practice often
  - an N-gram P(w |  $w_1 w_2 w_3 .... w_{k-1}$ )
  - a skipgram P(w |  $w_1 * w_3$ )





```
P("let's go to the party") =
P("let's | "<BOS>") *
P("go | "<BOS> let's" ) *
P("to" | ""<BOS> let's go") *
P("the" | "<BOS> let's go to") *
P("party" | "<BOS> let's go to the") *
P("<EOS>" | "<BOS> let's go to the party")
```





```
P("let's go to the party") =
P("let's | "<BOS>") *
P("go | "<BOS> let's" ) *
P("to" | ""<BOS> let's go") *
P("the" | let's go to") *
P("party" | "go to the") *
P("<EOS>" | to the party")
```





```
P("let's go to the party") =
P("let's | "<ROS ")
P("go | "<BOS> let's" ) *
P("to" | ""<BOS> let's go") *
P("the" | let's go to") *
P("party" | "go to the") *
P("<EOS>" | to the party")
```

by cutting down to 4-grams





## Unobserved word sequences

- Not all possible sequences of words may have occurred in the training material.
- Assigning unobserved sequences a probability equal to 0.0 will prohibit recognition of these sequences in a test.
- Therefore: assign these unobserved sequences a small probability by a clever estimate based on the text data in the training corpus





## Unobserved word sequences

Ways to assign probabilities to nonobserved word sequences:

- Smoothing: assigning the non-observed sequences (events) a small probability.
- Discounting: Estimate frequency of unobserved trigram by using the frequency of observed bigram of last two words.
   (Unobserved 4grams → 3grams, unobserved 3grams → bigrams, etc.)
- Clever methods available: e.g. "Kneser-Ney", "Good-Turing".





## bigram, ARPA file

```
\data\
ngram 1=<num 1-grams>
ngram 2=<num 2-ngrams>
\1-grams:
P(!ENTER)
              !ENTER B(!ENTER)
P(W1)
                      B(W1)
               W1
P(W2)
               W2
                      B(W2)
P(!EXIT)
               !EXIT
                     B(!EXIT)
\2-grams:
P(W1 | !ENTER) !ENTER W1
P(W2 | !ENTER) !ENTER W2
P(W1 | W1)
               W1
                       W1
P(W2 | W1)
               W1
                      W2
P(W1 | W2)
               W2
                       W1
P(!EXIT | W1)
               W1
                       !EXIT
P(!EXIT | W2)
               W2
                       EXIT
\end\
```





## How to create an LM for ASR

#### See for a detailed discussion

https://www.sciencedirect.com/topics/computer-science/language-modeling

#### There are many python tools to create LMs - e.g.

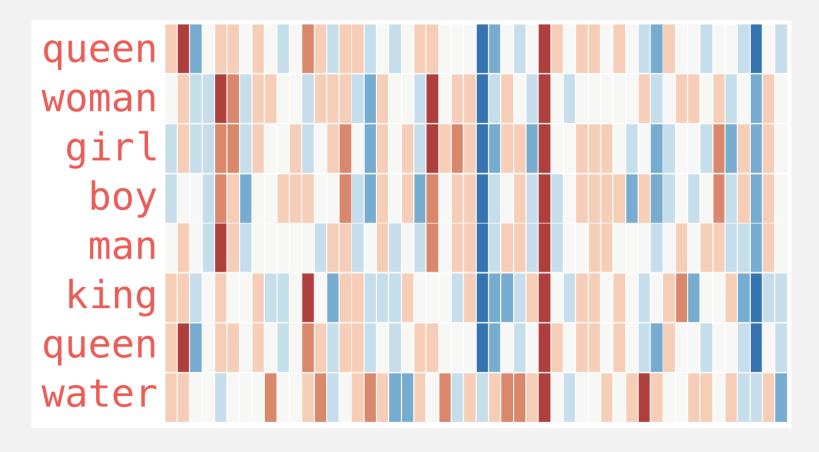
https://medium.com/analytics-vidhya/a-comprehensive-guide-to-build-your-own-language-model-in-python-5141b3917d6d

Steps	Examples
Raw	<a href="https://www.nc.nlm.new.edu.com/">httml ccdes&gt; Between A, B and C, there is a 3-way interaction, Mr. Bilmes reported.</a>
Clean up	Between A, B and C, there is a 3-way interaction, Mr. Bilmes reported.
Tokenisation	between A B and C here is a 3 way interaction Mr. Bilmes reported
Rewrites	between A B and C there is a three way interaction mister Bilmes reported
LM-building	output: arpa file, NN-LM





## Word embeddings, word2vec

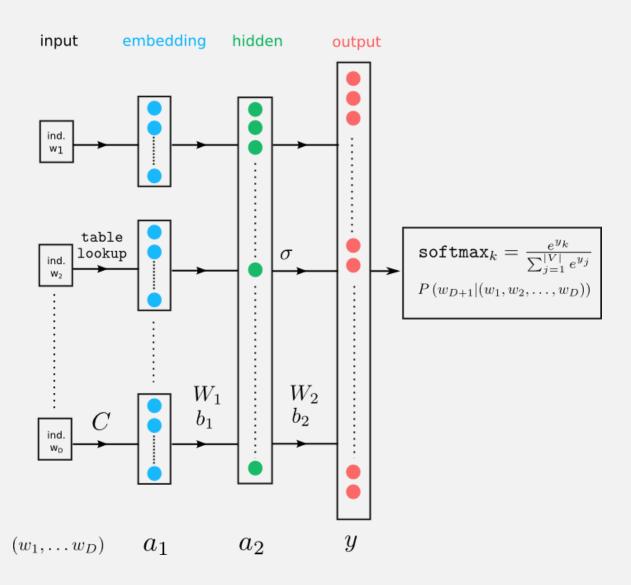


From: lajammar.github.io (retrieved 2020)



## NN language model

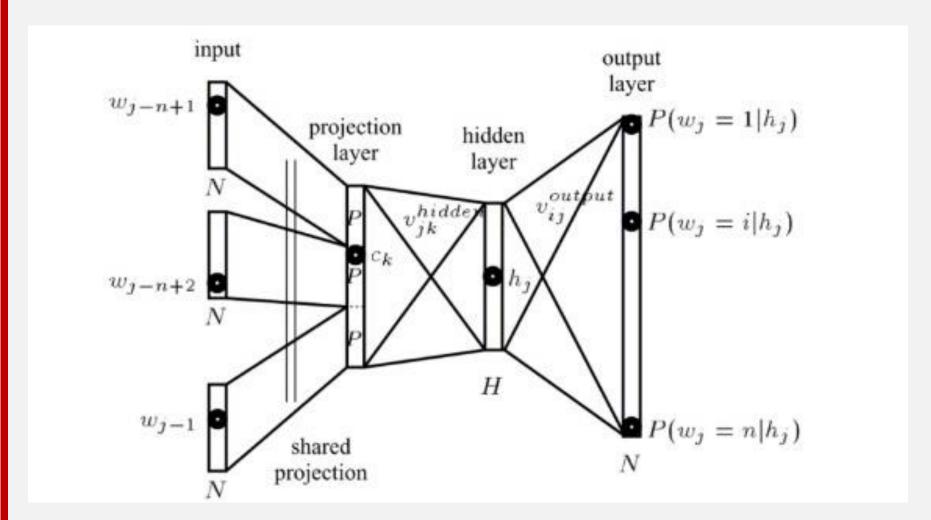








## (R)NN-LM







### Use of data

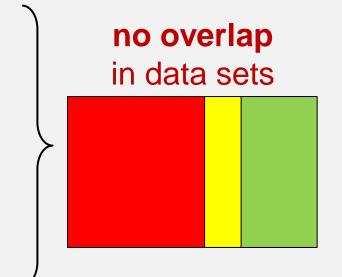
- Test data = training data
  - Weakness: very poor generalization power of model
  - Sometimes useful
- Separate training and test data (held out data)
  - Already better. Test on new unseen data.





### Better use of data

- training of the recognizer using a training set
- "tuning" of the recognizer using a development set
- testing of the recognizer using a held-out test set



- more advanced schemes:
  - Rotating schemes, N-fold cross validation schemes
  - Adversarial training and test methods



# Underfitting, overfitting, regularization



- An ASR system contains many parameters (usually millions)
- Models with lots of parameters can easily overfit to training data
- This overfitting might hamper generalization (the quality of a model on new, unseen data)
- Remedy: Regularization, by e.g. simplicity, sparsity, dropout, early stopping, Akaike Information Criterion (AIC), manifolds, etc.





## Evaluation of ASR quality

Usually based on a reference transcription

Expressed in terms of insertions, deletions, and substitutions

Example:

WER = 
$$\frac{\text{#del} + \text{#sub} + \text{#ins}}{N} \times 100\%$$
  
Wacc =  $\frac{\text{#corr} - \text{#del} - \text{#sub}}{N} \times 100\%$ 

N = number of words in reference





#### Three types of errors (at the phoneme level):

- 1. Deletion
- 2. Substitution
- 3. Insertion

#### **Phone Error Rate (PER)**

$$PER = \frac{dels + subs + ins}{N} \times 100\%$$

N = total number of phones in reference transcription





## Many variants possible

phonetic transcription: At bov@n d@n

reference transcription: z A t I o v @ n d @

- Option 1: consider hypothesis and reference sequences without taking into account the segmentation over time 

  Levenshtein
- Option 2: take time duration into account and construct a phone-phone confusion matrix frame by frame

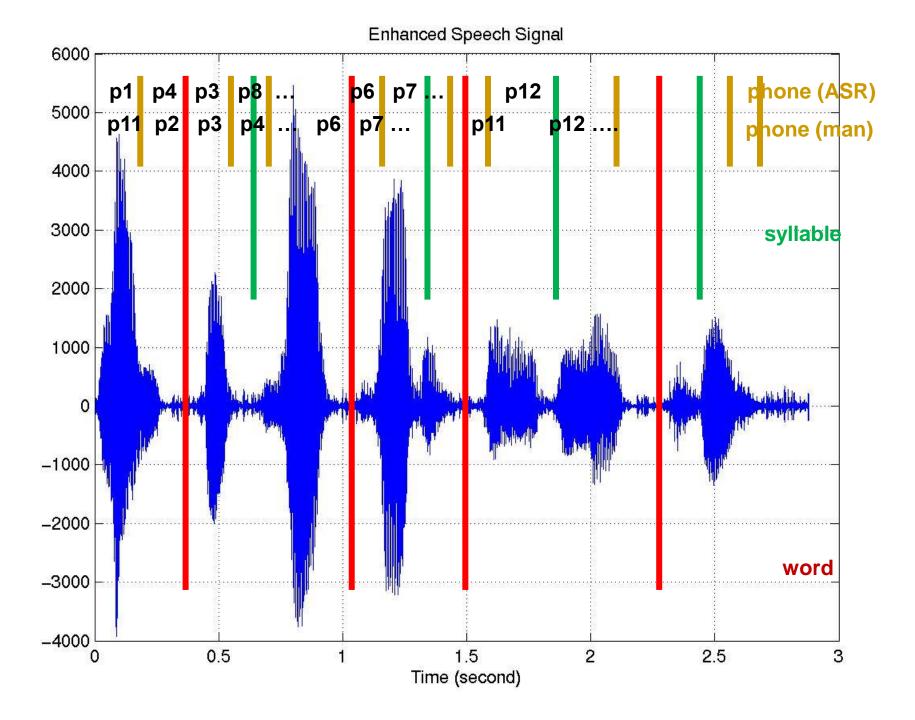




## Is WER a good measure?

- Usability of WER depends on the task and the performance
- For some applications, 70% word accuracy is OK; for other ones 98% is just not enough

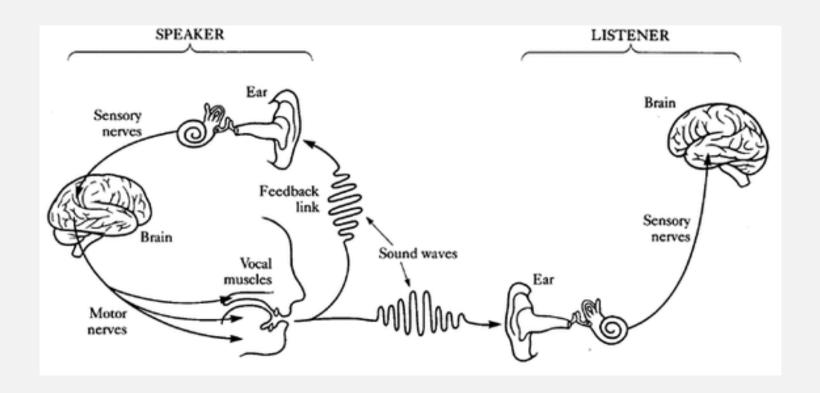
- WER still decreases (on average)
- ASR might soon become supra human





## What about humans?

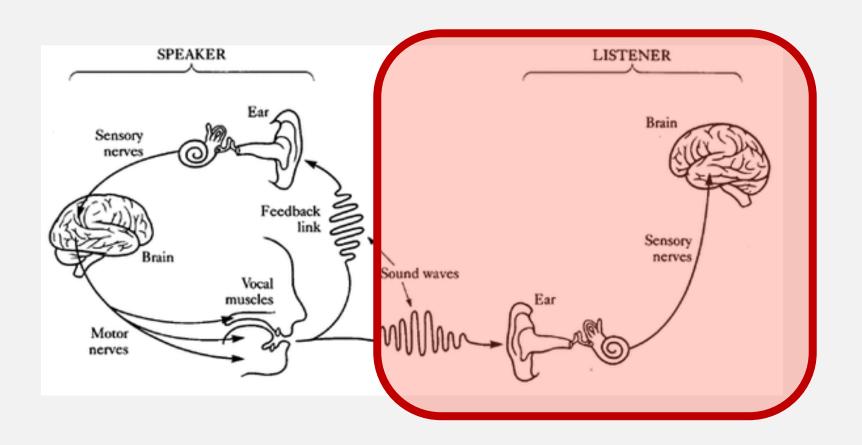




**Speaker-listener loop** 









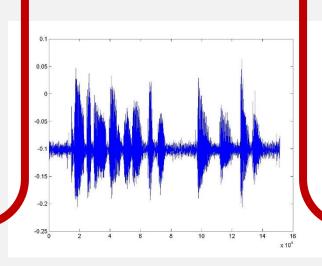


#### speaker

#### Concept

Grammar Mental lexicon

Muscle activation, articulation



#### listener

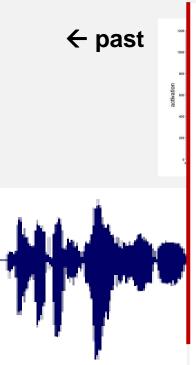
#### **Concept**

Grammar Mental lexicon

Auditory pathway



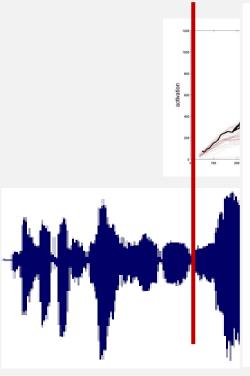




future →

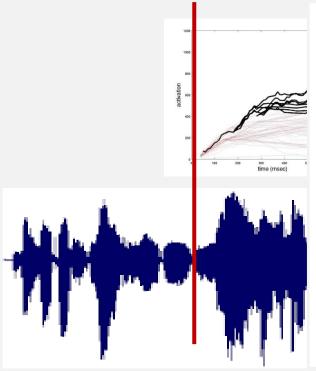






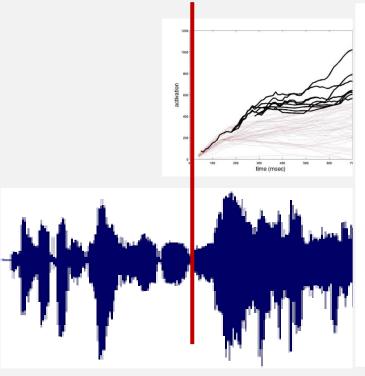






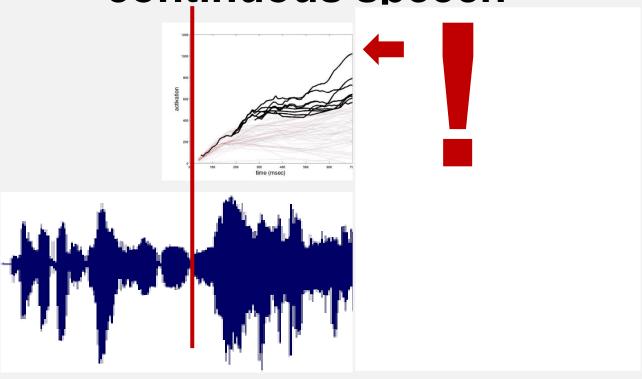








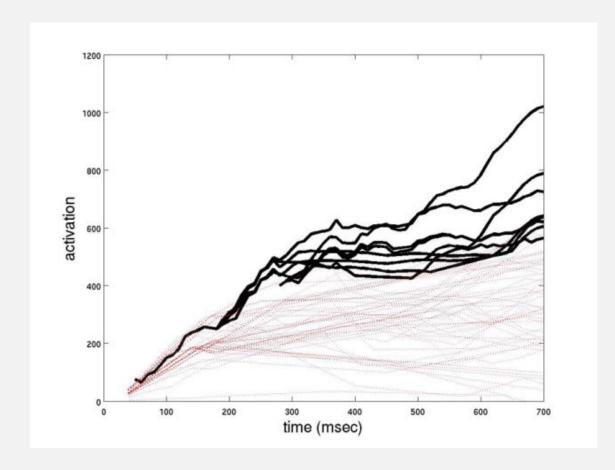






### Activation





Words compete in the listener's head

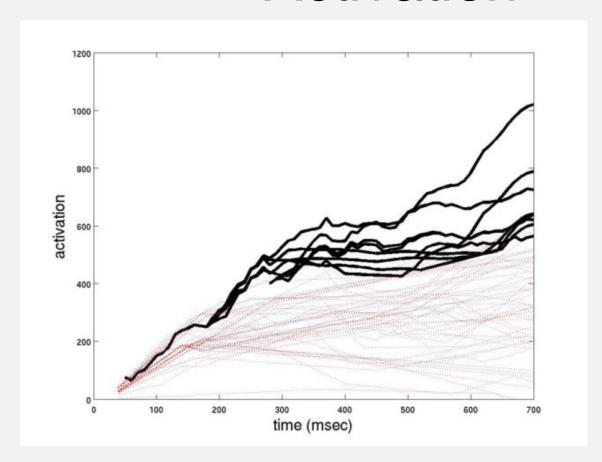
You have to select a winner 3 times a second

±60000 (±100) options per word





## Activation



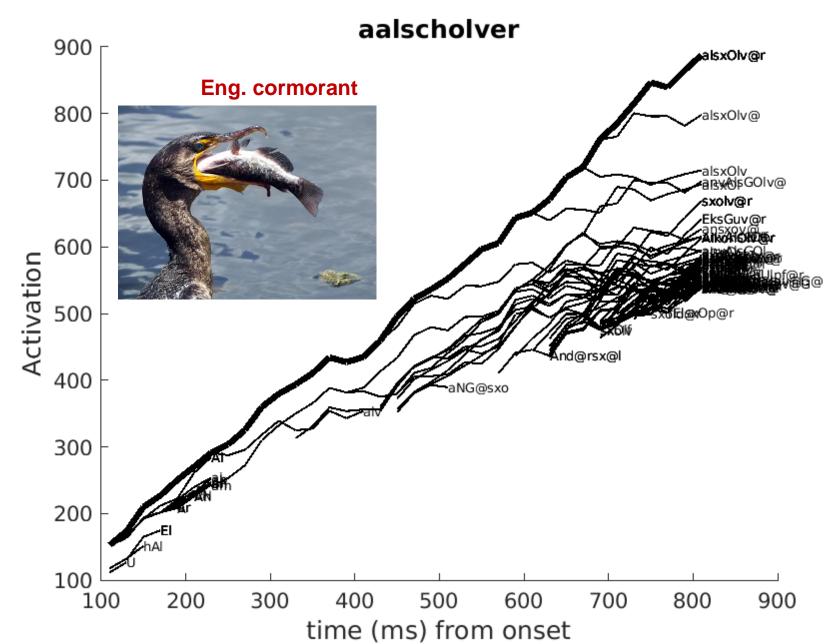
Activation =  $log(P(signal|word)) + \lambda log P(word)$ 

Example here: duration 0.68 sec.

Shown: top 200 candidates, recomputed each 10 ms

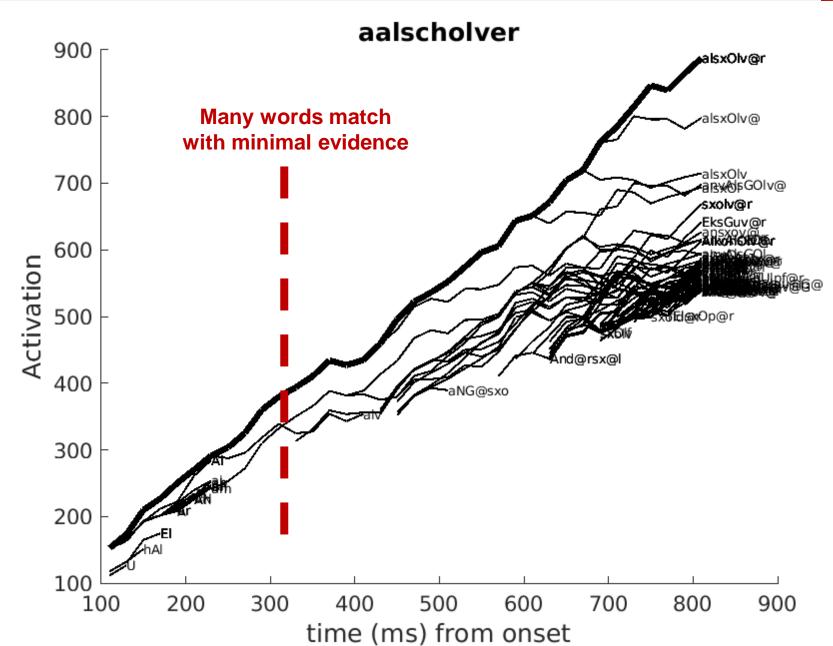






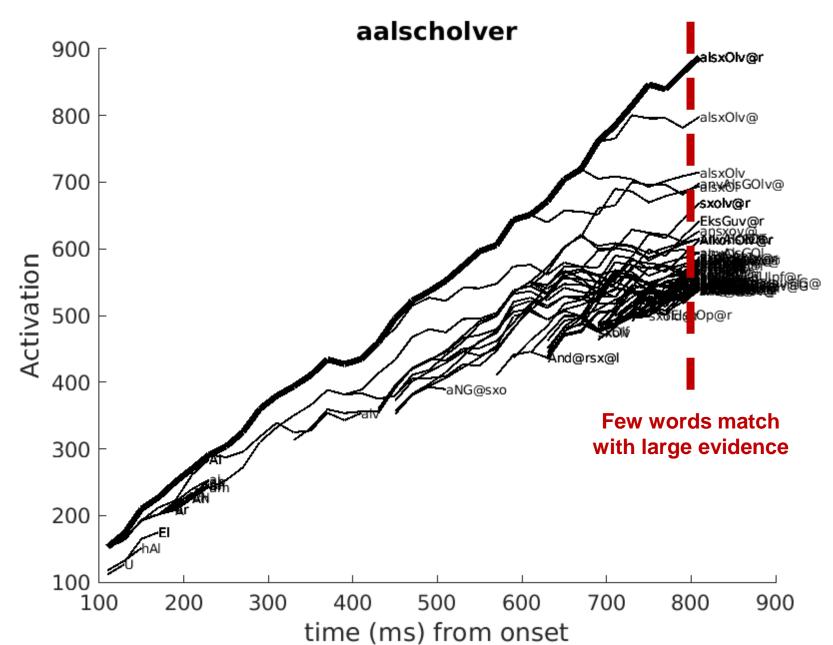






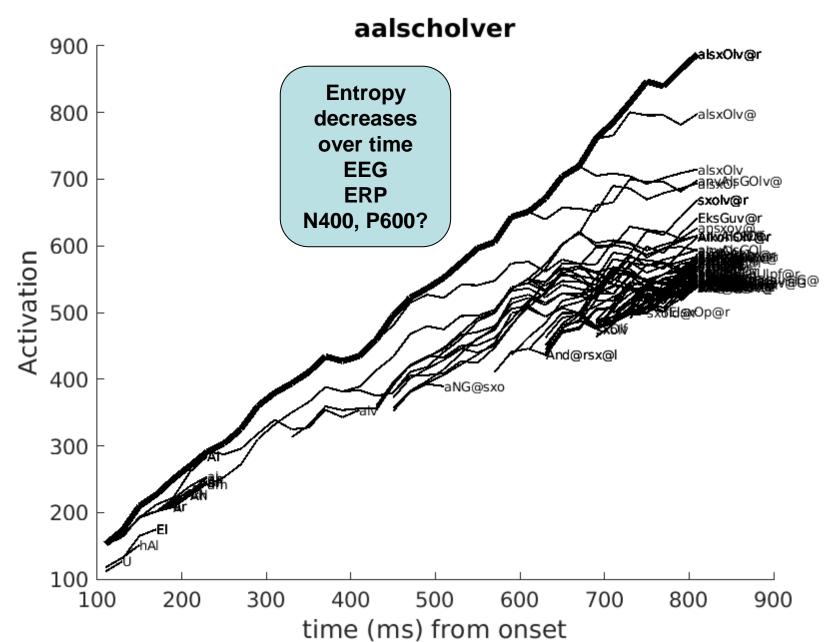










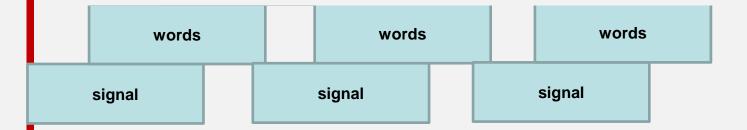




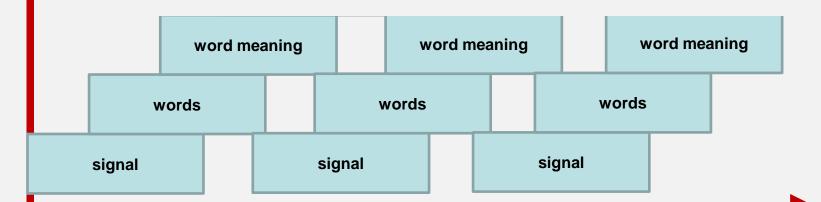
## Speech signal, words, meaning

signal signal signal

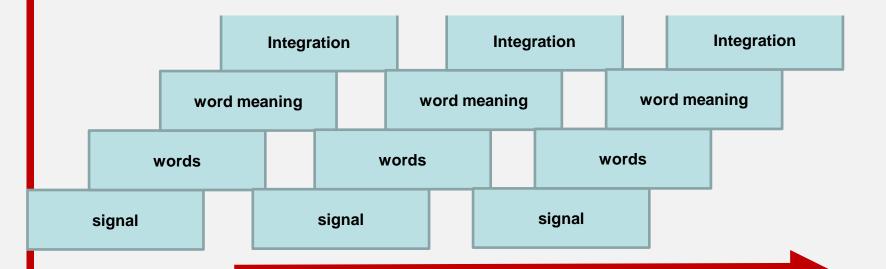




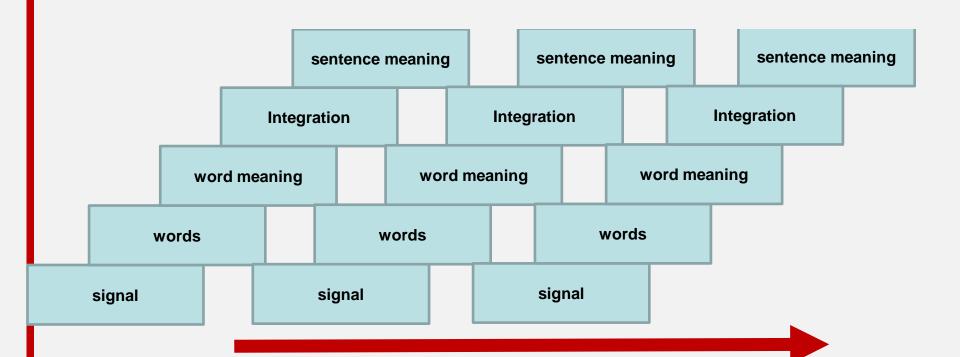




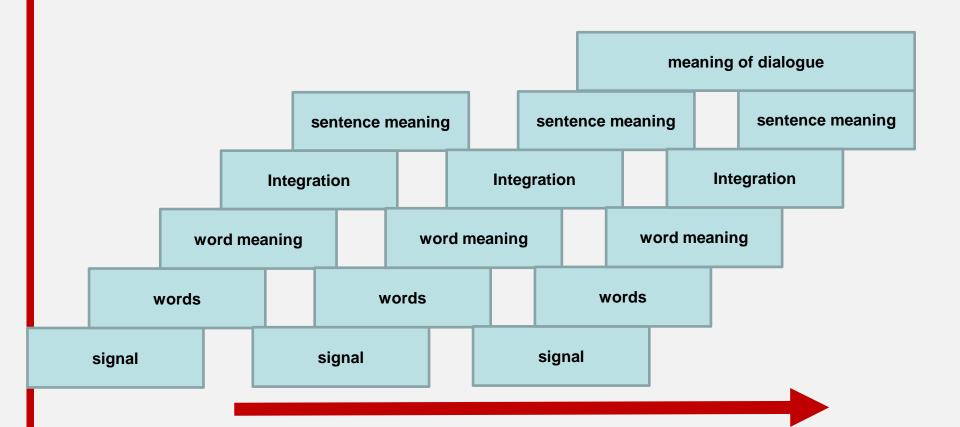




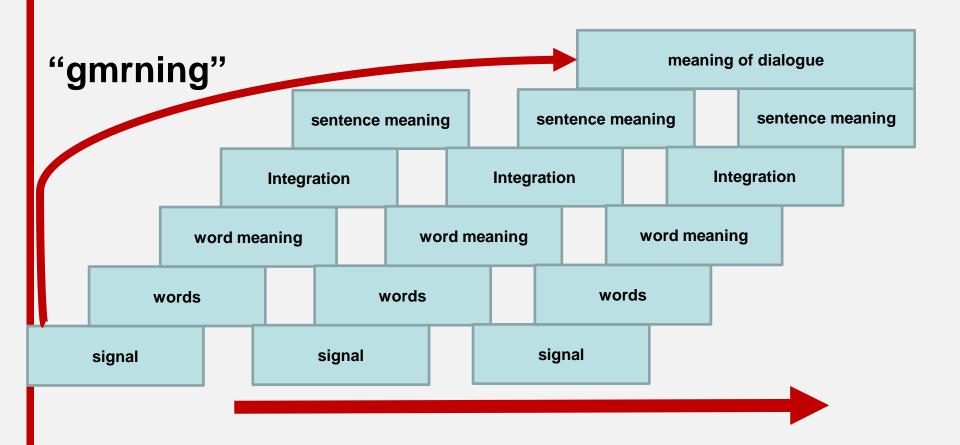








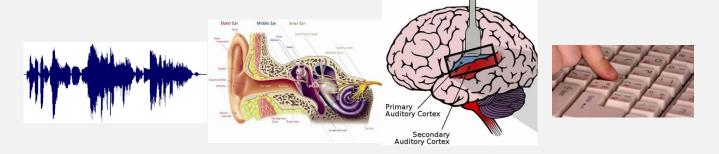






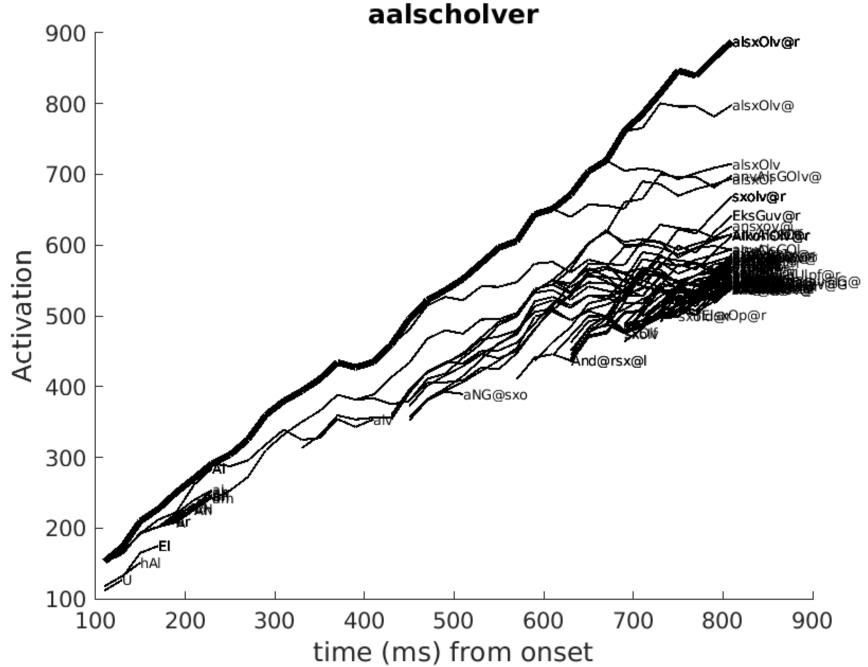


# Psycholinguistic experiment (e.g., lexical decision)

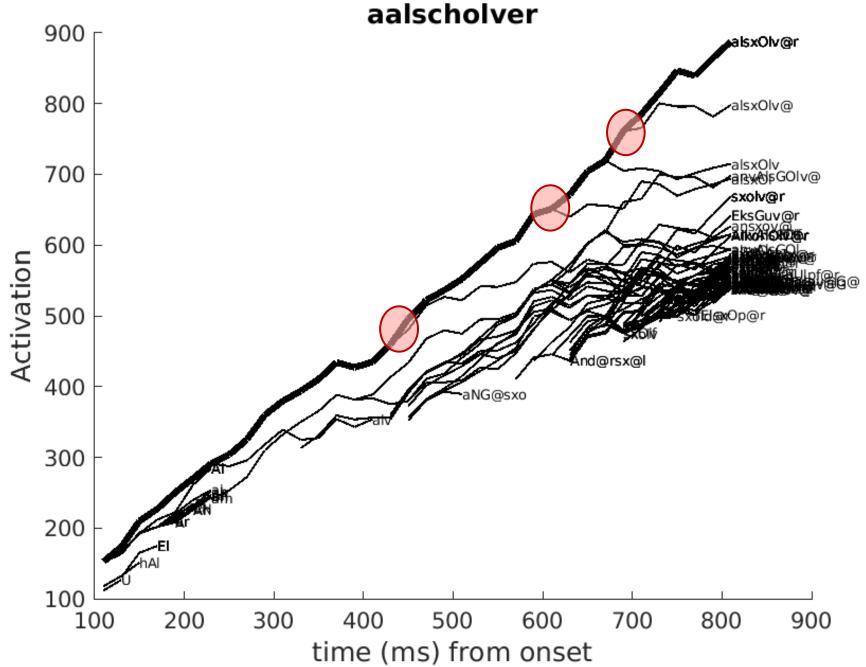


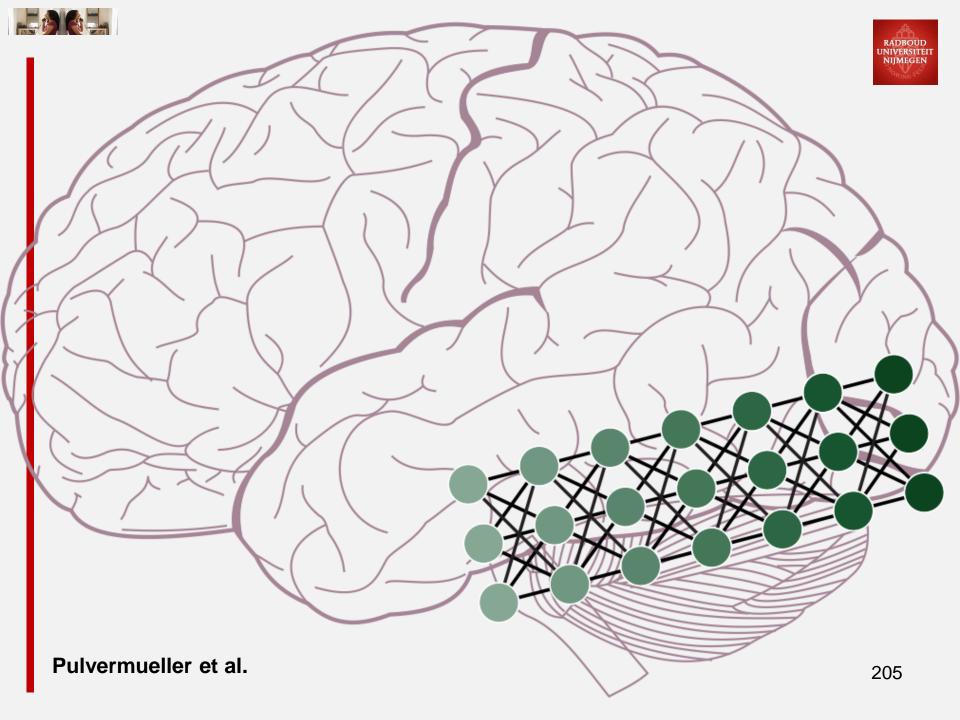
word activation, word competition decision execution reaction time (may be 1 sec!)















## Different ways to use ASR

- Free recognition
  - Wide search space
  - Broad topic
- Constrained recognition
  - Topic-based
  - Closed questions
- Forced alignment (wav + text)
  - Constrained search space
- Key word spotting
- Deep speech analysis





## Why DNN are useful: view 1

- DNNs may discover structure in data sets because subsequent layers ignore more and more details that are irrelevant for correctly predicting output labels
- Increasingly abstract representations emerge by cascading multiple (nonlinear) transformations
  - most convincing in image classification tasks.





#### View 2

- focus on the role of DNNs to find optimal representations, in particular in the sense of features.
- learning of optimal representations can be achieved if the network is able to disentangle the underlying explanatory factors hidden in the observed data.





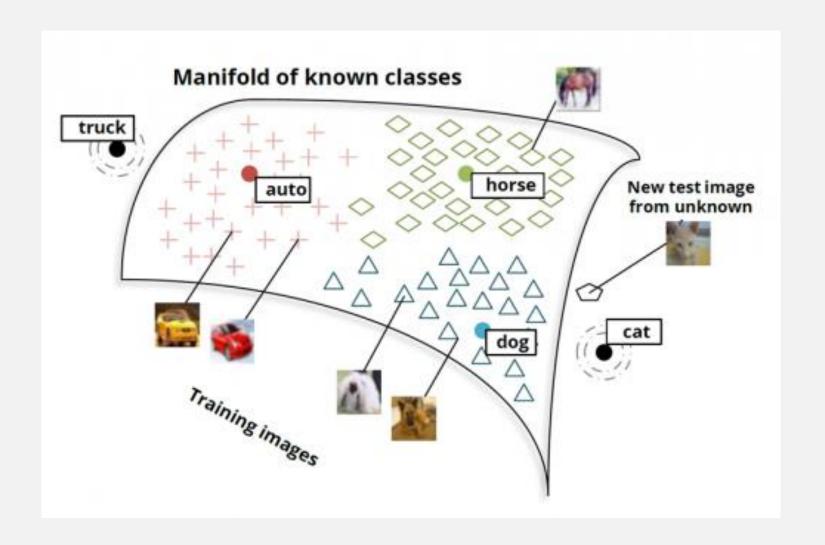
#### View 3

- more geometrically inspired interpretation of a DNN is based on the manifold assumption
  - the application of manifold learning methods on speech signals is (also) based on the relatively slow ballistic movements of articulators
- directions tangent to the manifold are well preserved while directions orthogonal to the manifolds aren't





### A manifold







#### View 4

- A fourth approach is more theoretical and analyzes DNNs on an 'information plane' using 'information bottleneck'
- Any DNN can be characterized by the mutual information between a hidden layer and the input and output variables, as a function of hidden layer depth
- bifurcation points of the information bottleneck trade-off





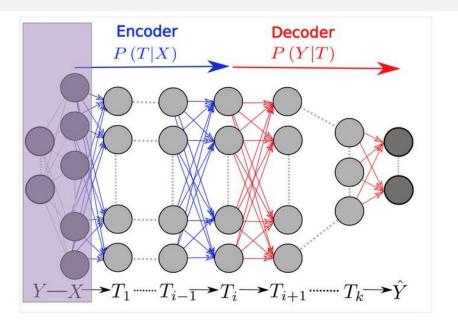


Figure 1: The DNN layers form a Markov chain of successive internal representations of the input layer X. Any representation of the input, T, is defined through an encoder, P(T|X), and a decoder  $P(\hat{Y}|T)$ , and can be quantified by its *information plane* coordinates:  $I_X = I(X;T)$  and  $I_Y = I(T;Y)$ . The Information Bottleneck bound characterizes the optimal representations, which maximally compress the input X, for a given mutual information on the desired output Y. After training, the network receives an input X, and successively processes it through the layers, which form a Markov chain, to the predicted output  $\hat{Y}$ .  $I(Y;\hat{Y})/I(X;Y)$  quantifies how much of the relevant information is captured by the network.





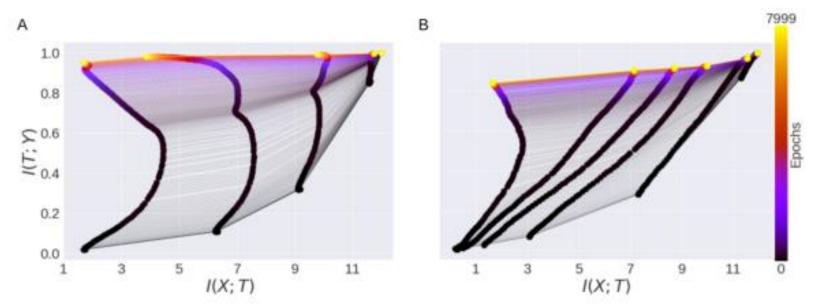


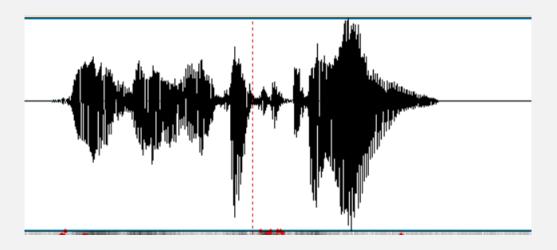
Figure 1: Information plane dynamics and neural nonlinearities. (A) Replication of Shwartz-Ziv & Tishby (2017) for a network with tanh nonlinearities (except for the final layer which contains sigmoidal neurons). The x-axis plots information between each layer and the input, while the y-axis plots information between each layer and the output. The color scale indicates training time in epochs. Each of the six layers produces a curve in the information plane with the input layer at far right, output layer at the far left. Different layers at the same epoch are connected by fine lines. (B) Information plane dynamics with ReLU nonlinearities (except for the final layer of 2 sigmoidal neurons). Here no compression phase is visible in the ReLU layers. For learning curves of both networks, see Appendix A



## Brand new approaches in ASR

- Focus on learning
  - Start with empty lexicon, and a large speech corpus along with images

'hey there look at the co-o-ow'











#### Given:

wav₁ + image₁

 $wav_2 + image_2$ 

 $wav_{\kappa} + image_{\kappa}$ 

#### Textual supervision for visually grounded spoken language understanding

Bertrand Higy\*, Desmond Elliott, Grzegorz Chrupala

\*Corresponding author for this work

Cognitive Science & Al

Research output: Chapter in Book/Report/Conference proceeding > Conference contribution > Scientific > peer-review

## How to dynamically build a lexicon?

- Lexicon contains links audio stretches (no spelling, no symbols yet) → image
- (Also an option for experiment)



